

# Intel Image Classification

The Intel Image Classification dataset is a collection of labeled images categorized into 6 types: **buildings, forest, glacier, mountain, sea, and street**. This dataset is intended for image classification tasks, providing a variety of real-world environmental categories. It can be used to train machine learning models for visual recognition tasks, helping in the development of models for image classification, especially in areas like environmental monitoring or geographical mapping.

## Required installations and libraries

Can be installed using the command `! pip install` below or by executing the installation with the `requirements.txt` file in the terminal using the command `pip install -r requirements.txt`.

```
# ! pip install numpy
# ! pip install pandas
# ! pip install matplotlib
# ! pip install tensorflow

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import json
import tensorflow as tf
from collections import Counter
from tensorflow import keras
from tensorflow.keras.utils import image_dataset_from_directory
from tensorflow.keras.utils import split_dataset
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
from tensorflow.keras import regularizers
from tensorflow.keras.applications import MobileNet
from tensorflow.keras.models import load_model
```

## Auxiliary code/functions

```
def counting_iterator(dataset_iterator, dataset_name,
    item_type='images'):
    """
    Counts the number of items (images or batches) in a TensorFlow
    dataset iterator.
    Prints the result as a message without returning anything.
```

```

Input:
    dataset_iterator: tf.data.Dataset
        A TensorFlow dataset iterator containing the data.
    dataset_name: str
        Name of the dataset (e.g., 'train', 'test', 'validation').
    item_type: str, optional
        Specifies whether we are counting individual 'images' or
        'batches'. Defaults to 'images'.

Output:
    None
    """
    total_items = sum(1 for _ in dataset_iterator)

    print(f'The {dataset_name} dataset has {total_items}
    {item_type}.')

def count_items_by_label(dataset, labels_names):
    """
    Counts the number of items per label in a dataset.

    Input:
        dataset: tf.data.Dataset
            A TensorFlow dataset iterator containing the data.
        labels_names: list
            A list of class labels corresponding to the indices in the
            dataset.

    Output:
        label_counting: dict
            A dictionary with the class index as keys and the number
            of items as values.
        """

    label_counting = Counter()

    for _, label in dataset:
        class_idx = np.argmax(label.numpy()) # converting the
        # tensor label to a numpy array and obtaining the index with argmax
        label_counting[class_idx] += 1 # label_counting will
        # be a dictionary with label_idx as keys and number of images as values

    for index, count in label_counting.items():
        print(f'{labels_names[index]}: {count} images.')

    return label_counting

def show_images(dataset, class_names, num_images=16, rows=4, cols=4):
    """
    Adapted from:

```

[https://www.tensorflow.org/tutorials/load\\_data/images#visualize\\_the\\_data](https://www.tensorflow.org/tutorials/load_data/images#visualize_the_data) in order to

display a grid of 16 images from a dataset along with their corresponding class labels in a 4x4 grid.

*Input:*

`dataset: tf.data.Dataset`

A TensorFlow dataset that yields image-label pairs.

`class_names: list`

A list of class names corresponding to the one-hot encoded labels. The index of each class name in this list corresponds to the class label.

`num_images : int, optional`

The number of images to display. Default is 16.

`rows : int, optional`

The number of rows in the image grid. Default is 4.

`cols : int, optional`

The number of columns in the image grid. Default is 4.

*Output:*

None

Simply displays the images and their corresponding labels in a grid layout.

"""

`plt.figure(figsize=(10, 10))`

`for i, (image, label) in enumerate(dataset.take(num_images)):`

`if i >= num_images:`

`break`

`plt.subplot(rows, cols, i+1)`

`plt.imshow(image.numpy().astype("uint8"))`

`plt.axis('off')`

`class_name = class_names[label.numpy().argmax()]`

`plt.title(f'{class_name}')`

`plt.show()`

`def plot_accuracy_f1_loss(history, epochs):`

"""

*Adapted from:*

<https://www.tensorflow.org/tutorials/images/classification> to plot the training and validation accuracy, F1 score, and loss metrics during training.

*Input:*

`history: keras.callbacks.History`

*The history object returned by the model.fit() function.*

*epochs: int*

*The number of epochs used during training. This is used to determine the range for the x-axis in the plots.*

*Output:*

*None*

*Display a grid of 3 plots: Training and Validation Accuracy, F1 Score and Loss.*

```
"""
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

f1_score = history.history['f1_score']
val_f1_score = history.history['val_f1_score']

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs_range = range(epochs)

plt.figure(figsize=(10,6))

plt.subplot(1, 3, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Accuracy')

plt.subplot(1, 3, 2)
plt.plot(epochs_range, f1_score, label='Training F1 Score')
plt.plot(epochs_range, val_f1_score, label='Validation F1 Score')
plt.legend(loc='lower right')
plt.title('F1 Score')

plt.subplot(1, 3, 3)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Loss')

plt.show()

def save_model_log(model_name, history, epochs, file_path='log.json'):
    """
    Save the best performance metrics of the model to a JSON log file.
```

*Input:*

*model\_name: str*

*The name of the model being logged.*

*history: keras.callbacks.History*  
*The history object returned by the model.fit() function, containing the training metrics.*

*epochs: int*  
*The number of epochs used during training. This is used to track the training process.*

*file\_path: str, optional (default='log.json')*  
*The file path where the log data will be saved. If the file exists, the data will be appended.*

*Output:*  
*None*  
*The function saves the best metrics to a JSON file, without returning any value.*

```
"""  
  
best_training_loss = min(history.history['loss'])  
best_training_accuracy = max(history.history['accuracy'])  
best_training_f1_score = max(history.history['f1_score'])  
  
best_val_loss = min(history.history['val_loss'])  
best_val_accuracy = max(history.history['val_accuracy'])  
best_val_f1_score = max(history.history['val_f1_score'])  
  
log_data = {  
    'model_name': model_name,  
    'epochs': epochs,  
    'configurations': {  
        'optimizer': 'Adam',  
        'loss_function': 'CategoricalCrossentropy  
(from_logits=True)',  
        'metrics': ['Accuracy', 'F1 Score']  
    },  
    'best_metrics': {  
        'training': {  
            'loss': best_training_loss,  
            'accuracy': best_training_accuracy,  
            'f1_score': best_training_f1_score  
        },  
        'validation': {  
            'loss': best_val_loss,  
            'accuracy': best_val_accuracy,  
            'f1_score': best_val_f1_score  
        }  
    }  
}
```

```

try:
    with open(file_path, 'a') as f:
        json.dump(log_data, f, indent=4)
        f.write('\n')

except Exception as e:
    print(f'An error occurred: {e}')
    with open(file_path, 'w') as f:
        json.dump(log_data, f, indent=4)

```

## Load the dataset

I have downloaded the dataset from the Kaggle repository available at the following link: <https://www.kaggle.com/datasets/puneet6060/intel-image-classification/data>.

According to the dataset description, it contains approximately **25 000 images of size 150x150, distributed across 6 categories**: {0: Buildings, 1: Forest, 2: Glacier, 3: Mountain, 4: Sea, 5: Street}.

The data is already split into 3 subsets: training (~14 000 images), testing (~3 000 images), and prediction (~7 000 images).

Since the purpose of this project is to rapidly iterate through the training process, I will only upload one split, specifically the training set, which will be further divided later on.

```
dataset_directory = '../Project_DeepLearning/input/seg_train'
```

To properly load the dataset, I referred to the documentation available at [https://keras.io/api/data\\_loading/image/#image-data-loading](https://keras.io/api/data_loading/image/#image-data-loading).

This will return a `tf.data.Dataset` object, where images are represented as tensors of shape (batch\_size, image\_height, image\_width, num\_channels) and labels are tensors of shape (num\_labels,), corresponding to one-hot encoded vectors.

```

from tensorflow import keras
from tensorflow.keras.utils import image_dataset_from_directory

dataset = keras.utils.image_dataset_from_directory(
    dataset_directory,
    labels='inferred',
    label_mode='categorical',
    batch_size=None,
    image_size=(150, 150),
    shuffle=True,
    # Labels will be named as the
    # folders' names
    # To get the labels as one-hot
    # encoded vectors
    # Image Size as per the
    # documentation of the dataset
)

```

```

    seed=42
)
Found 14034 files belonging to 6 classes.
dataset
<_PrefetchDataset element_spec=(TensorSpec(shape=(150, 150, 3),
dtype=tf.float32, name=None), TensorSpec(shape=(6,), dtype=tf.float32,
name=None))>
for image, label in dataset:
    print(f'Image shape: {image.shape}')
    print(f'Label shape: {label.shape}')
    print(label)
    break
Image shape: (150, 150, 3)
Label shape: (6,)
tf.Tensor([0. 0. 0. 0. 1. 0.], shape=(6,), dtype=float32)

```

The shape of the image is (150, 150, 3), which means the images are in color with **3 RGB channels**.

```

dataset.class_names
['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']

```

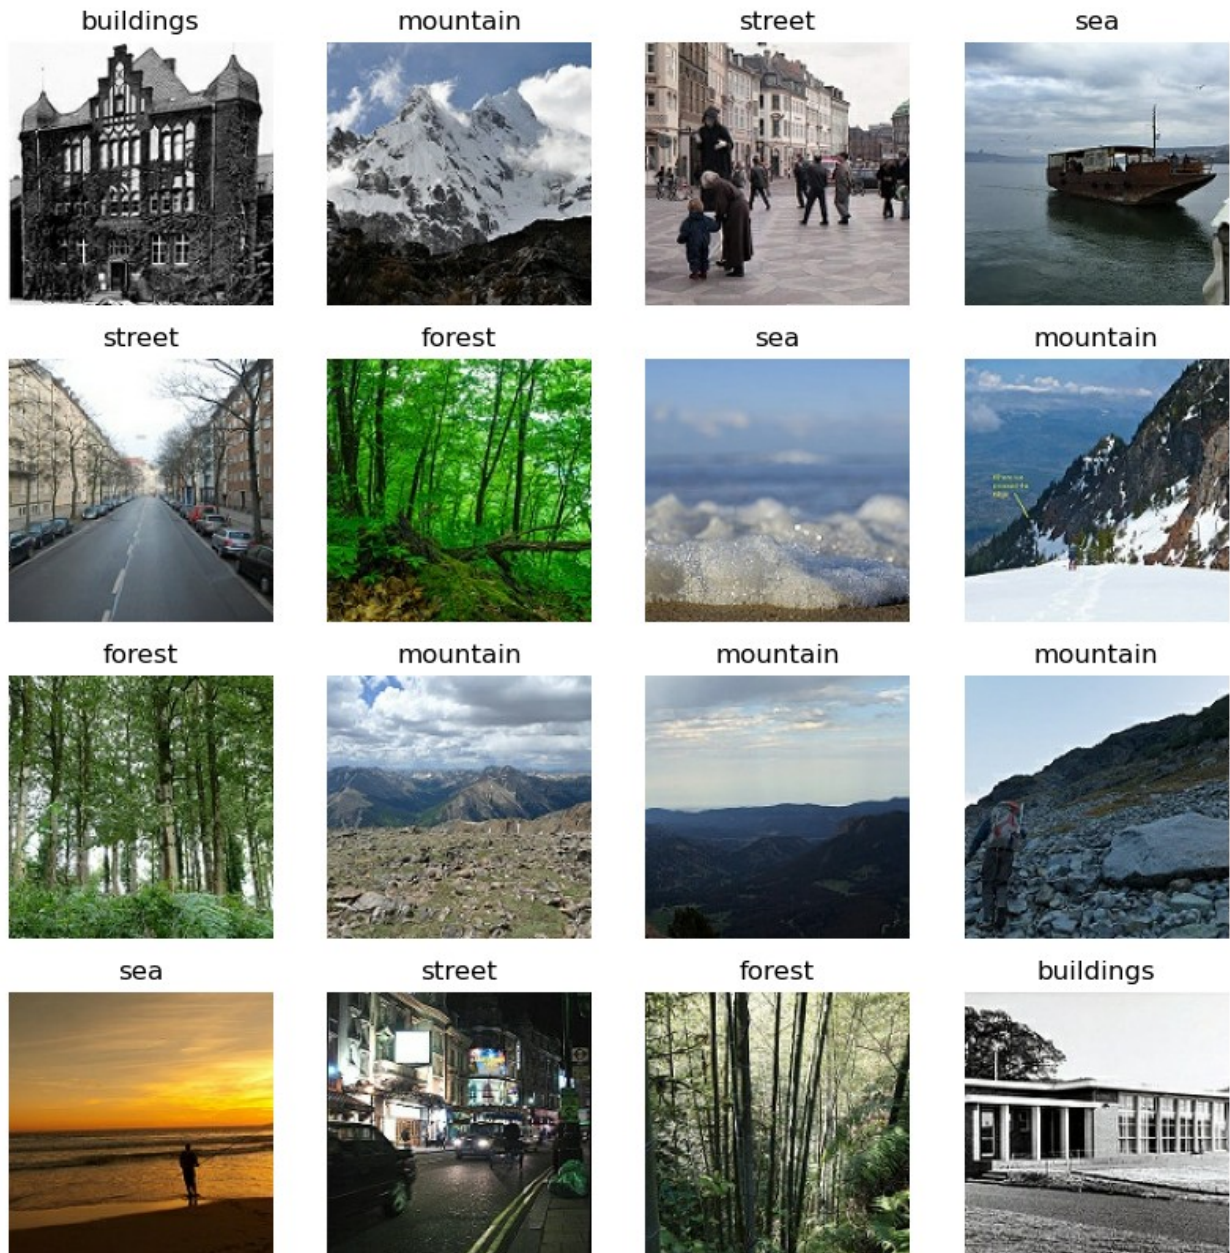
## Visualize the data

```

labels_names = dataset.class_names
print(labels_names)
['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
show_images(dataset, labels_names)

```





## Explore the dataset

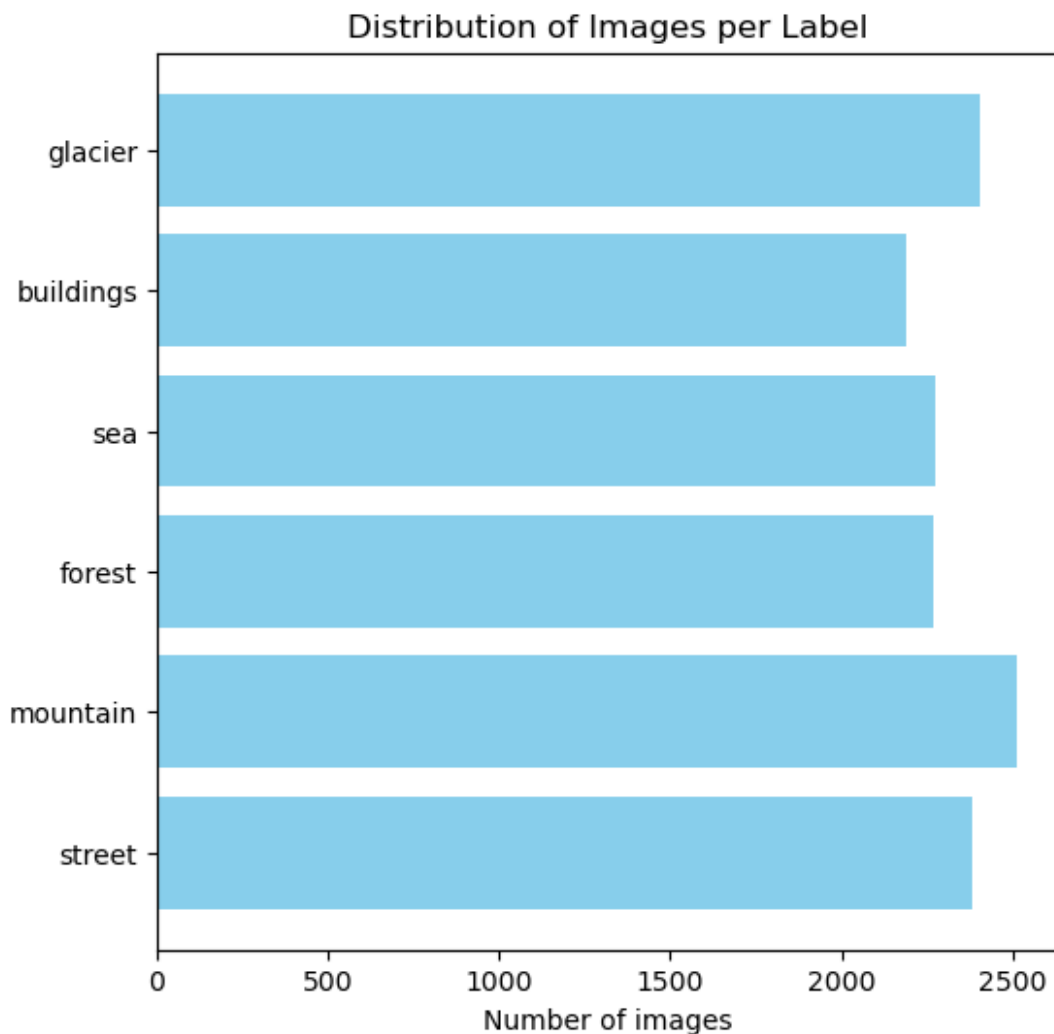
Upon loading the dataset, we observe that it contains 14 034 images across 6 classes. However, will verify whether the classes are balanced or not.

```
n_images_per_label = count_items_by_label(dataset, labels_names)
```

```
street: 2382 images.  
mountain: 2512 images.  
forest: 2271 images.  
sea: 2274 images.
```



```
buildings: 2191 images.  
glacier: 2404 images.  
  
# Visualize graphically the distribution of images per class  
  
classes = [labels_names[i] for i in n_images_per_label.keys()]  
counts = [n_images_per_label[i] for i in n_images_per_label.keys()]  
  
plt.figure(figsize=(6, 6))  
plt.barh(classes, counts, color='skyblue')  
plt.xlabel('Number of images')  
plt.title('Distribution of Images per Label')  
plt.show()
```



# Transform the dataset

## Train, test, validation splits

Since the purpose of this project is to build a Neural Network from scratch and test different architectures, I will split the dataset into training, validation, and testing sets with fewer examples: 3 000 for training, 1 500 for validation, and 600 for testing.

For this, I have used the `split_dataset` function as referenced in the following link:  
[https://keras.io/api/utils/python\\_utils/#splitdataset-function](https://keras.io/api/utils/python_utils/#splitdataset-function).

```
from tensorflow.keras.utils import split_dataset

train_size = 3000
val_size = 1500
test_size = 600

# Dividing the dataset into the training dataset with 3000 images, and the rest

train_ds, rest = tf.keras.utils.split_dataset(
    dataset,
    left_size=train_size,
    seed=42
)

# Now with the rest dataset from above, I will divide into validation and test datasets

val_ds, test_ds = tf.keras.utils.split_dataset(
    rest,
    left_size=val_size,
    right_size=test_size,
    seed=42
)

counting_iterator(train_ds, 'train'),
counting_iterator(val_ds, 'test'),
counting_iterator(test_ds, 'validation')

The train dataset has 3000 images.
The test dataset has 1500 images.
The validation dataset has 600 images.

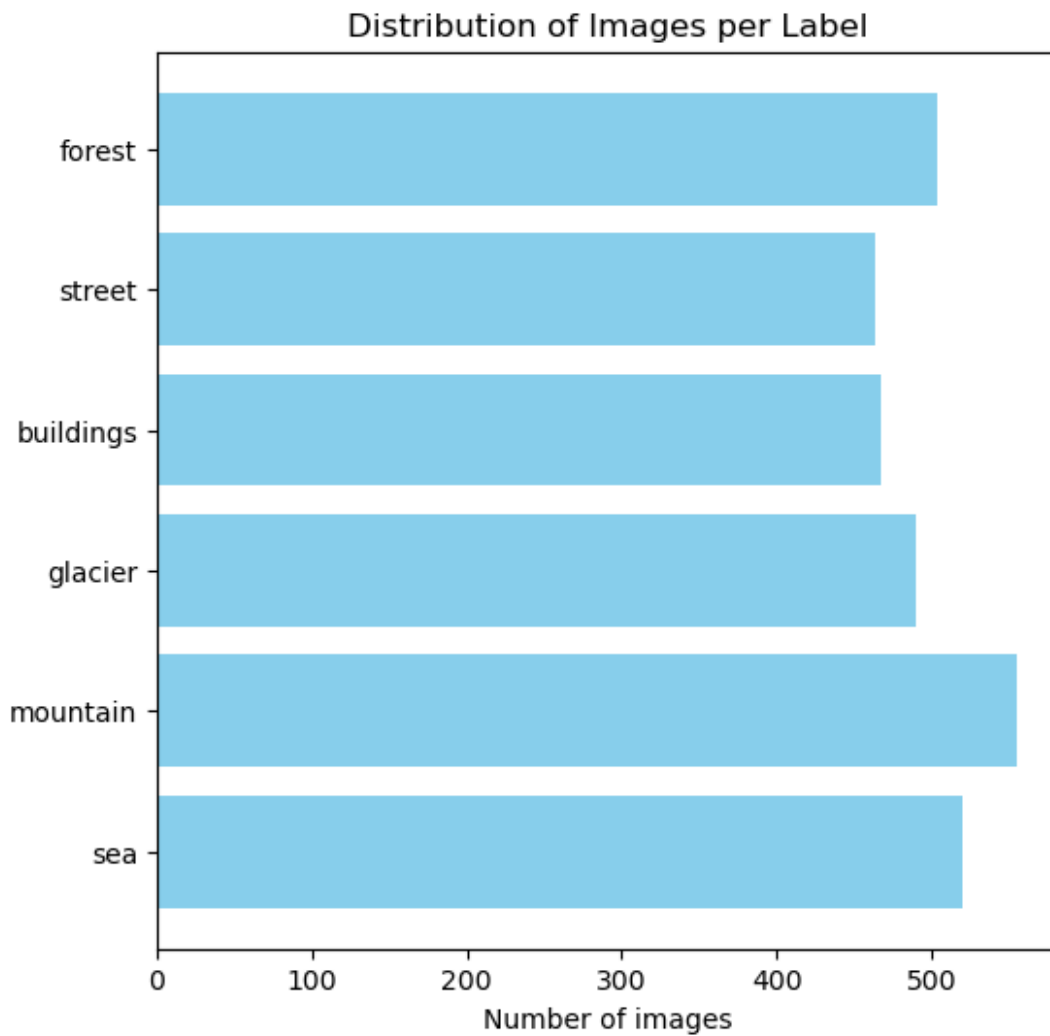
label_n_train = count_items_by_label(train_ds, labels_names)

# Visualize graphically the distribution of images per class

classes_train = [labels_names[i] for i in label_n_train.keys()]
counts_train = [label_n_train[i] for i in label_n_train.keys()]
```

```
plt.figure(figsize=(6, 6))
plt.barh(classes_train, counts_train, color='skyblue')
plt.xlabel('Number of images')
plt.title('Distribution of Images per Label')
plt.show()
```

```
sea: 520 images.
mountain: 555 images.
glacier: 490 images.
buildings: 467 images.
street: 464 images.
forest: 504 images.
```



```
label_n_val = count_items_by_label(val_ds, labels_names)

mountain: 260 images.
sea: 233 images.
```

```
street: 267 images.
glacier: 273 images.
forest: 247 images.
buildings: 220 images.

label_n_test = count_items_by_label(test_ds, labels_names)

glacier: 101 images.
forest: 83 images.
street: 95 images.
mountain: 120 images.
sea: 109 images.
buildings: 92 images.
```

## Normalize Data

```
def normalize_image(image, label):
    return tf.cast(image, tf.float32)/255.0, label

normalize_train_ds = train_ds.map(normalize_image)
normalize_val_ds = val_ds.map(normalize_image)
normalize_test_ds = test_ds.map(normalize_image)

# Checking if values were normalized:

for image, label in normalize_train_ds.take(1):
    print(image.shape)
    print("Min:", tf.reduce_min(image).numpy())
    print("Max:", tf.reduce_max(image).numpy())

(150, 150, 3)
Min: 0.0
Max: 1.0
```

## Define Batch Size

```
train_dataset = normalize_train_ds.batch(256)
test_dataset = normalize_test_ds.batch(256)
val_dataset = normalize_val_ds.batch(256)

counting_iterator(train_dataset, 'train', 'batches')
counting_iterator(test_dataset, 'test', 'batches')
counting_iterator(val_dataset, 'validation', 'batches')

The train dataset has 12 batches.
The test dataset has 3 batches.
The validation dataset has 6 batches.
```

The datasets have now been preprocessed and transformed, and are ready for the training and validation process.

# Modelling

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
```

## Metrics

I chose **Accuracy** and **F1 Score** as evaluation metrics.

Accuracy provides a general sense of the model's performance, but it can be misleading in imbalanced datasets. Although the Intel Image Classification dataset is not imbalanced, I'll use F1 Score to maintain robustness and prevent any potential issues in this regard. Additionally, the F1 Score harmonizes both precision and recall, ensuring a more comprehensive assessment of the model, rather than relying solely on accuracy.

## Baseline model

I will start by building a model using only Dense layers, to establish a baseline for comparing the performance of Dense-based models with Convolutional-based ones.

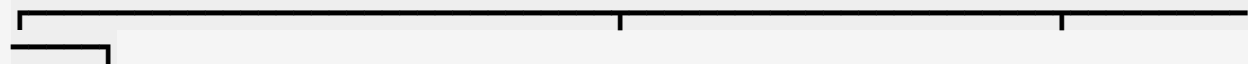
```
num_classes = len(labels_names)

baseline_model = Sequential([
    layers.Input(shape=(150, 150, 3)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(64, activation='relu'),
    layers.Dense(32, activation='relu'),
    layers.Dense(num_classes)
])

baseline_model.compile(
    optimizer='adam',
    loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
    metrics=[
        'accuracy',
        tf.keras.metrics.F1Score(average='macro', name='f1_score')
    ]
)

baseline_model.summary()

Model: "sequential"
```



Layer (type) Param #	Output Shape
flatten (Flatten)	(None, 67500)
dense (Dense)	(None, 128)
dense_1 (Dense)	(None, 64)
dense_2 (Dense)	(None, 32)
dense_3 (Dense)	(None, 6)

Total params: 8,650,662 (33.00 MB)

Trainable params: 8,650,662 (33.00 MB)

Non-trainable params: 0 (0.00 B)

```
epochs=50
history = baseline_model.fit(
    train_dataset,
    validation_data=val_dataset,
    epochs=epochs
)
```

Epoch 1/50

12/12 ————— 2s 70ms/step - accuracy: 0.1625 - f1\_score: 0.1065 - loss: 11.9336 - val\_accuracy: 0.2660 - val\_f1\_score: 0.1374 - val\_loss: 3.8544

Epoch 2/50

12/12 ————— 1s 61ms/step - accuracy: 0.2159 - f1\_score: 0.1556 - loss: 3.2858 - val\_accuracy: 0.1520 - val\_f1\_score: 0.0554 - val\_loss: 2.3844

Epoch 3/50

12/12 ————— 1s 57ms/step - accuracy: 0.2065 - f1\_score: 0.1831 - loss: 2.1423 - val\_accuracy: 0.3233 - val\_f1\_score: 0.2528 - val\_loss: 1.6568

Epoch 4/50  
12/12 \_\_\_\_\_ 1s 57ms/step - accuracy: 0.3144 - f1\_score: 0.3053 - loss: 1.6614 - val\_accuracy: 0.4620 - val\_f1\_score: 0.4160 - val\_loss: 1.4634

Epoch 5/50  
12/12 \_\_\_\_\_ 1s 62ms/step - accuracy: 0.3954 - f1\_score: 0.3755 - loss: 1.5075 - val\_accuracy: 0.4360 - val\_f1\_score: 0.3942 - val\_loss: 1.4384

Epoch 6/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.4219 - f1\_score: 0.4012 - loss: 1.4800 - val\_accuracy: 0.4633 - val\_f1\_score: 0.4316 - val\_loss: 1.3726

Epoch 7/50  
12/12 \_\_\_\_\_ 1s 58ms/step - accuracy: 0.4494 - f1\_score: 0.4328 - loss: 1.4296 - val\_accuracy: 0.4667 - val\_f1\_score: 0.4366 - val\_loss: 1.3798

Epoch 8/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.4589 - f1\_score: 0.4449 - loss: 1.3880 - val\_accuracy: 0.4560 - val\_f1\_score: 0.4133 - val\_loss: 1.3693

Epoch 9/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.4874 - f1\_score: 0.4727 - loss: 1.3228 - val\_accuracy: 0.4713 - val\_f1\_score: 0.4338 - val\_loss: 1.3419

Epoch 10/50  
12/12 \_\_\_\_\_ 1s 62ms/step - accuracy: 0.4977 - f1\_score: 0.4870 - loss: 1.2974 - val\_accuracy: 0.4740 - val\_f1\_score: 0.4389 - val\_loss: 1.3302

Epoch 11/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.5142 - f1\_score: 0.5066 - loss: 1.2757 - val\_accuracy: 0.4973 - val\_f1\_score: 0.4547 - val\_loss: 1.3103

Epoch 12/50  
12/12 \_\_\_\_\_ 1s 62ms/step - accuracy: 0.5265 - f1\_score: 0.5171 - loss: 1.2538 - val\_accuracy: 0.4893 - val\_f1\_score: 0.4547 - val\_loss: 1.3103

Epoch 13/50  
12/12 \_\_\_\_\_ 1s 59ms/step - accuracy: 0.5254 - f1\_score: 0.5185 - loss: 1.2391 - val\_accuracy: 0.5013 - val\_f1\_score: 0.4553 - val\_loss: 1.2956

Epoch 14/50  
12/12 \_\_\_\_\_ 1s 62ms/step - accuracy: 0.5442 - f1\_score: 0.5368 - loss: 1.2171 - val\_accuracy: 0.4867 - val\_f1\_score: 0.4548 - val\_loss: 1.2966

Epoch 15/50  
12/12 \_\_\_\_\_ 1s 63ms/step - accuracy: 0.5452 - f1\_score: 0.5398 - loss: 1.2038 - val\_accuracy: 0.5100 - val\_f1\_score: 0.4675 - val\_loss: 1.2829

Epoch 16/50



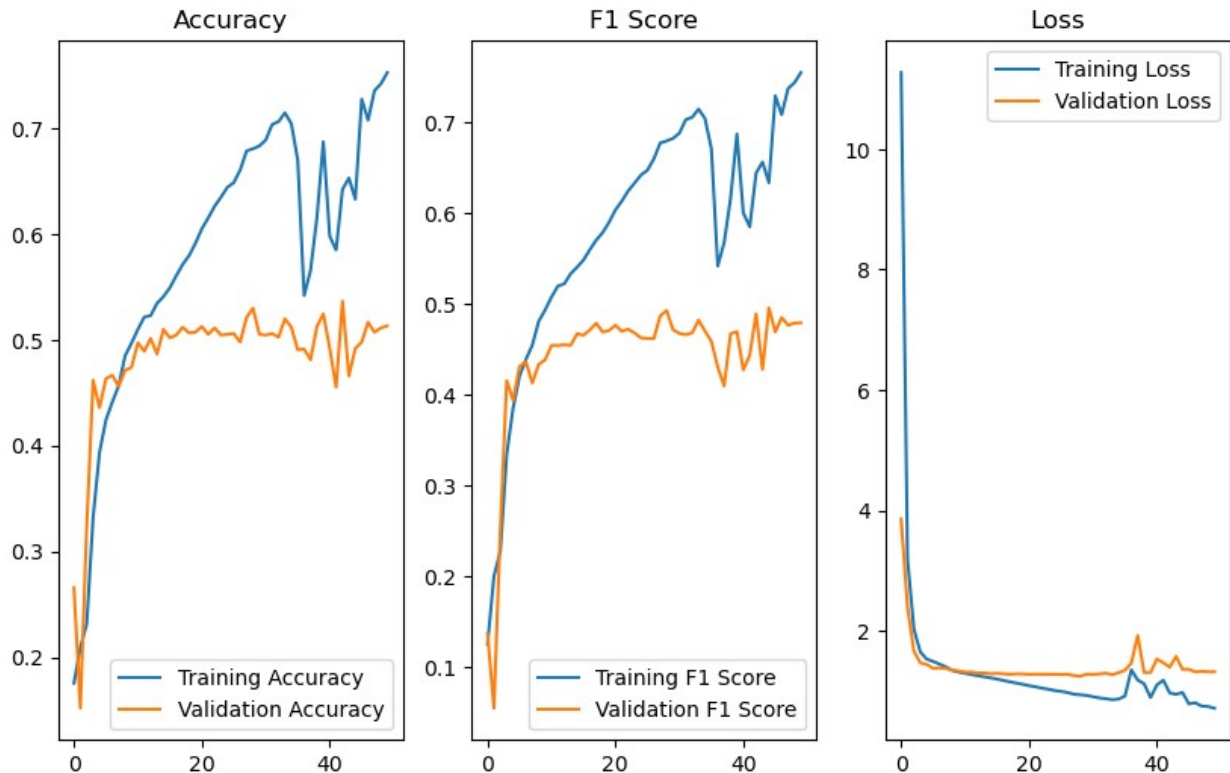
12/12 \_\_\_\_\_ 1s 66ms/step - accuracy: 0.5597 - f1\_score: 0.5543 - loss: 1.1834 - val\_accuracy: 0.5020 - val\_f1\_score: 0.4660 - val\_loss: 1.2891  
Epoch 17/50  
12/12 \_\_\_\_\_ 1s 61ms/step - accuracy: 0.5667 - f1\_score: 0.5604 - loss: 1.1606 - val\_accuracy: 0.5047 - val\_f1\_score: 0.4717 - val\_loss: 1.2826  
Epoch 18/50  
12/12 \_\_\_\_\_ 1s 69ms/step - accuracy: 0.5737 - f1\_score: 0.5680 - loss: 1.1411 - val\_accuracy: 0.5120 - val\_f1\_score: 0.4791 - val\_loss: 1.2747  
Epoch 19/50  
12/12 \_\_\_\_\_ 1s 67ms/step - accuracy: 0.5835 - f1\_score: 0.5788 - loss: 1.1196 - val\_accuracy: 0.5067 - val\_f1\_score: 0.4695 - val\_loss: 1.2804  
Epoch 20/50  
12/12 \_\_\_\_\_ 1s 64ms/step - accuracy: 0.5953 - f1\_score: 0.5896 - loss: 1.0996 - val\_accuracy: 0.5073 - val\_f1\_score: 0.4711 - val\_loss: 1.2805  
Epoch 21/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.6091 - f1\_score: 0.6038 - loss: 1.0775 - val\_accuracy: 0.5127 - val\_f1\_score: 0.4770 - val\_loss: 1.2721  
Epoch 22/50  
12/12 \_\_\_\_\_ 1s 59ms/step - accuracy: 0.6210 - f1\_score: 0.6160 - loss: 1.0595 - val\_accuracy: 0.5053 - val\_f1\_score: 0.4703 - val\_loss: 1.2760  
Epoch 23/50  
12/12 \_\_\_\_\_ 1s 61ms/step - accuracy: 0.6289 - f1\_score: 0.6257 - loss: 1.0371 - val\_accuracy: 0.5113 - val\_f1\_score: 0.4726 - val\_loss: 1.2738  
Epoch 24/50  
12/12 \_\_\_\_\_ 1s 62ms/step - accuracy: 0.6335 - f1\_score: 0.6309 - loss: 1.0230 - val\_accuracy: 0.5047 - val\_f1\_score: 0.4683 - val\_loss: 1.2735  
Epoch 25/50  
12/12 \_\_\_\_\_ 1s 61ms/step - accuracy: 0.6444 - f1\_score: 0.6420 - loss: 0.9998 - val\_accuracy: 0.5053 - val\_f1\_score: 0.4630 - val\_loss: 1.2734  
Epoch 26/50  
12/12 \_\_\_\_\_ 1s 61ms/step - accuracy: 0.6464 - f1\_score: 0.6437 - loss: 0.9891 - val\_accuracy: 0.5060 - val\_f1\_score: 0.4624 - val\_loss: 1.2713  
Epoch 27/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.6592 - f1\_score: 0.6574 - loss: 0.9718 - val\_accuracy: 0.4980 - val\_f1\_score: 0.4622 - val\_loss: 1.2751  
Epoch 28/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.6752 - f1\_score:

0.6739 - loss: 0.9443 - val\_accuracy: 0.5213 - val\_f1\_score: 0.4870 -  
val\_loss: 1.2522  
Epoch 29/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.6770 - f1\_score:  
0.6761 - loss: 0.9391 - val\_accuracy: 0.5300 - val\_f1\_score: 0.4932 -  
val\_loss: 1.2410  
Epoch 30/50  
12/12 \_\_\_\_\_ 1s 57ms/step - accuracy: 0.6814 - f1\_score:  
0.6808 - loss: 0.9268 - val\_accuracy: 0.5053 - val\_f1\_score: 0.4721 -  
val\_loss: 1.2756  
Epoch 31/50  
12/12 \_\_\_\_\_ 1s 59ms/step - accuracy: 0.6894 - f1\_score:  
0.6894 - loss: 0.9030 - val\_accuracy: 0.5047 - val\_f1\_score: 0.4679 -  
val\_loss: 1.2722  
Epoch 32/50  
12/12 \_\_\_\_\_ 1s 58ms/step - accuracy: 0.7023 - f1\_score:  
0.7021 - loss: 0.8847 - val\_accuracy: 0.5060 - val\_f1\_score: 0.4666 -  
val\_loss: 1.2796  
Epoch 33/50  
12/12 \_\_\_\_\_ 1s 57ms/step - accuracy: 0.7062 - f1\_score:  
0.7070 - loss: 0.8712 - val\_accuracy: 0.5027 - val\_f1\_score: 0.4684 -  
val\_loss: 1.2917  
Epoch 34/50  
12/12 \_\_\_\_\_ 1s 59ms/step - accuracy: 0.7141 - f1\_score:  
0.7152 - loss: 0.8548 - val\_accuracy: 0.5200 - val\_f1\_score: 0.4825 -  
val\_loss: 1.2679  
Epoch 35/50  
12/12 \_\_\_\_\_ 1s 59ms/step - accuracy: 0.7034 - f1\_score:  
0.7035 - loss: 0.8646 - val\_accuracy: 0.5120 - val\_f1\_score: 0.4706 -  
val\_loss: 1.2998  
Epoch 36/50  
12/12 \_\_\_\_\_ 1s 59ms/step - accuracy: 0.6803 - f1\_score:  
0.6801 - loss: 0.8982 - val\_accuracy: 0.4907 - val\_f1\_score: 0.4592 -  
val\_loss: 1.3410  
Epoch 37/50  
12/12 \_\_\_\_\_ 1s 60ms/step - accuracy: 0.5760 - f1\_score:  
0.5740 - loss: 1.2274 - val\_accuracy: 0.4913 - val\_f1\_score: 0.4310 -  
val\_loss: 1.4623  
Epoch 38/50  
12/12 \_\_\_\_\_ 1s 58ms/step - accuracy: 0.5768 - f1\_score:  
0.5742 - loss: 1.1152 - val\_accuracy: 0.4813 - val\_f1\_score: 0.4101 -  
val\_loss: 1.9187  
Epoch 39/50  
12/12 \_\_\_\_\_ 1s 57ms/step - accuracy: 0.5886 - f1\_score:  
0.5761 - loss: 1.2690 - val\_accuracy: 0.5127 - val\_f1\_score: 0.4675 -  
val\_loss: 1.3026  
Epoch 40/50  
12/12 \_\_\_\_\_ 1s 59ms/step - accuracy: 0.6978 - f1\_score:  
0.6955 - loss: 0.8742 - val\_accuracy: 0.5247 - val\_f1\_score: 0.4699 -

```
val_loss: 1.3036
Epoch 41/50
12/12 _____ 1s 61ms/step - accuracy: 0.6340 - f1_score:
0.6317 - loss: 1.0024 - val_accuracy: 0.4920 - val_f1_score: 0.4277 -
val_loss: 1.5286
Epoch 42/50
12/12 _____ 1s 59ms/step - accuracy: 0.5688 - f1_score:
0.5632 - loss: 1.2323 - val_accuracy: 0.4553 - val_f1_score: 0.4439 -
val_loss: 1.4658
Epoch 43/50
12/12 _____ 1s 59ms/step - accuracy: 0.6213 - f1_score:
0.6243 - loss: 1.0202 - val_accuracy: 0.5367 - val_f1_score: 0.4894 -
val_loss: 1.3920
Epoch 44/50
12/12 _____ 1s 58ms/step - accuracy: 0.6610 - f1_score:
0.6564 - loss: 0.9305 - val_accuracy: 0.4660 - val_f1_score: 0.4283 -
val_loss: 1.5694
Epoch 45/50
12/12 _____ 1s 57ms/step - accuracy: 0.5975 - f1_score:
0.5952 - loss: 1.0496 - val_accuracy: 0.4920 - val_f1_score: 0.4961 -
val_loss: 1.3564
Epoch 46/50
12/12 _____ 1s 59ms/step - accuracy: 0.7197 - f1_score:
0.7225 - loss: 0.7936 - val_accuracy: 0.4980 - val_f1_score: 0.4696 -
val_loss: 1.3572
Epoch 47/50
12/12 _____ 1s 61ms/step - accuracy: 0.6902 - f1_score:
0.6891 - loss: 0.8229 - val_accuracy: 0.5167 - val_f1_score: 0.4853 -
val_loss: 1.3140
Epoch 48/50
12/12 _____ 1s 60ms/step - accuracy: 0.7263 - f1_score:
0.7236 - loss: 0.7625 - val_accuracy: 0.5073 - val_f1_score: 0.4769 -
val_loss: 1.3221
Epoch 49/50
12/12 _____ 1s 62ms/step - accuracy: 0.7270 - f1_score:
0.7250 - loss: 0.7584 - val_accuracy: 0.5113 - val_f1_score: 0.4793 -
val_loss: 1.3157
Epoch 50/50
12/12 _____ 1s 61ms/step - accuracy: 0.7417 - f1_score:
0.7394 - loss: 0.7316 - val_accuracy: 0.5133 - val_f1_score: 0.4796 -
val_loss: 1.3156

save_model_log('baseline_model', history, epochs)

plot_accuracy_f1_loss(history, epochs)
```



It is now evident why Dense layers are not ideal for extracting meaningful features from images. Not only are the metrics generally underwhelming, but the validation metrics also quickly stagnated around 50%, indicating poor performance.

## Models with Convolutional Layers

I will now build a neural network using Convolutional layers, which are known to be more effective for image classification. Therefore, I expect it to outperform the baseline model.

I'll start with a shallow architecture and gradually increase its depth, exploring techniques such as Padding, Max Pooling and regularization techniques along the way.

```
num_classes = len(labels_names)

cnn_base = Sequential([
    layers.Input(shape=(150, 150, 3)),
    layers.Conv2D(16, 3, activation='relu'),
    layers.Conv2D(32, 3, activation='relu'),
    layers.Flatten(),
    layers.Dense(32, activation='relu'),
    layers.Dense(num_classes)
])

cnn_base.compile(
    optimizer='adam',
    loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
```

```

        metrics=[
            'accuracy',
            tf.keras.metrics.F1Score(average='macro', name='f1_score')
        ]
    )

```

```

cnn_base.summary()

```

```

Model: "sequential_1"

```

Layer (type) Param #	Output Shape	
conv2d (Conv2D) 448	(None, 148, 148, 16)	
conv2d_1 (Conv2D) 4,640	(None, 146, 146, 32)	
flatten_1 (Flatten) 0	(None, 682112)	
dense_4 (Dense) 21,827,616	(None, 32)	
dense_5 (Dense) 198	(None, 6)	

```

Total params: 21,832,902 (83.29 MB)

```

```

Trainable params: 21,832,902 (83.29 MB)

```

```

Non-trainable params: 0 (0.00 B)

```

```

epochs = 20

```

```

history_cnn_base = cnn_base.fit(
    train_dataset,
    validation_data=val_dataset,
    epochs=epochs
)

```

Epoch 1/20  
12/12 \_\_\_\_\_ 10s 690ms/step - accuracy: 0.2096 -  
f1\_score: 0.1778 - loss: 10.6220 - val\_accuracy: 0.3613 -  
val\_f1\_score: 0.2513 - val\_loss: 1.6546

Epoch 2/20  
12/12 \_\_\_\_\_ 8s 635ms/step - accuracy: 0.3789 -  
f1\_score: 0.2606 - loss: 1.5414 - val\_accuracy: 0.3553 - val\_f1\_score:  
0.2663 - val\_loss: 1.3623

Epoch 3/20  
12/12 \_\_\_\_\_ 8s 639ms/step - accuracy: 0.5026 -  
f1\_score: 0.4726 - loss: 1.1855 - val\_accuracy: 0.5680 - val\_f1\_score:  
0.5217 - val\_loss: 1.1002

Epoch 4/20  
12/12 \_\_\_\_\_ 8s 663ms/step - accuracy: 0.6313 -  
f1\_score: 0.6072 - loss: 0.9301 - val\_accuracy: 0.5820 - val\_f1\_score:  
0.5558 - val\_loss: 1.0787

Epoch 5/20  
12/12 \_\_\_\_\_ 8s 648ms/step - accuracy: 0.7292 -  
f1\_score: 0.7263 - loss: 0.7023 - val\_accuracy: 0.6300 - val\_f1\_score:  
0.5997 - val\_loss: 1.0358

Epoch 6/20  
12/12 \_\_\_\_\_ 8s 644ms/step - accuracy: 0.8396 -  
f1\_score: 0.8414 - loss: 0.5056 - val\_accuracy: 0.6493 - val\_f1\_score:  
0.6211 - val\_loss: 1.0107

Epoch 7/20  
12/12 \_\_\_\_\_ 8s 638ms/step - accuracy: 0.9014 -  
f1\_score: 0.9035 - loss: 0.3522 - val\_accuracy: 0.6273 - val\_f1\_score:  
0.5908 - val\_loss: 1.0737

Epoch 8/20  
12/12 \_\_\_\_\_ 8s 647ms/step - accuracy: 0.9428 -  
f1\_score: 0.9437 - loss: 0.2477 - val\_accuracy: 0.6180 - val\_f1\_score:  
0.5970 - val\_loss: 1.0949

Epoch 9/20  
12/12 \_\_\_\_\_ 8s 649ms/step - accuracy: 0.9543 -  
f1\_score: 0.9558 - loss: 0.1976 - val\_accuracy: 0.6633 - val\_f1\_score:  
0.6415 - val\_loss: 1.0091

Epoch 10/20  
12/12 \_\_\_\_\_ 8s 646ms/step - accuracy: 0.9693 -  
f1\_score: 0.9702 - loss: 0.1685 - val\_accuracy: 0.6153 - val\_f1\_score:  
0.6059 - val\_loss: 1.0679

Epoch 11/20  
12/12 \_\_\_\_\_ 8s 649ms/step - accuracy: 0.9445 -  
f1\_score: 0.9446 - loss: 0.2045 - val\_accuracy: 0.5460 - val\_f1\_score:  
0.5358 - val\_loss: 1.3667

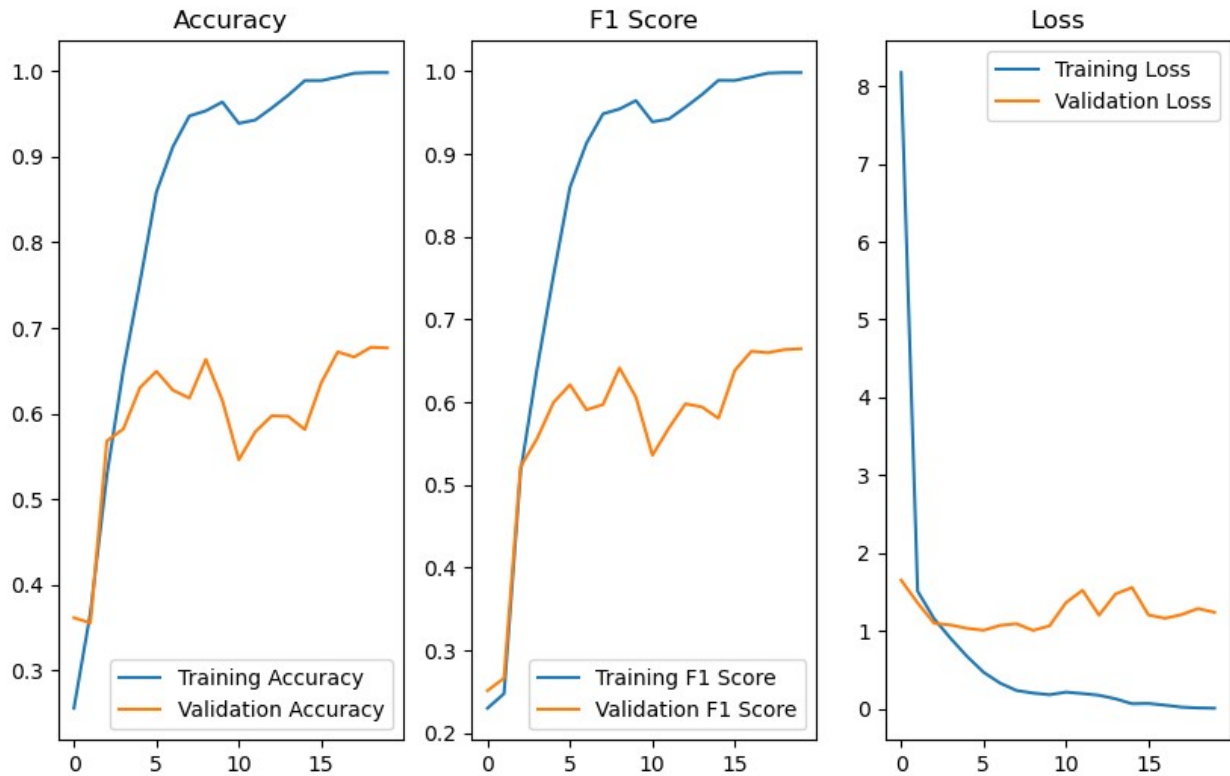
Epoch 12/20  
12/12 \_\_\_\_\_ 8s 652ms/step - accuracy: 0.9323 -  
f1\_score: 0.9302 - loss: 0.2281 - val\_accuracy: 0.5787 - val\_f1\_score:  
0.5688 - val\_loss: 1.5244

Epoch 13/20  
12/12 \_\_\_\_\_ 8s 656ms/step - accuracy: 0.9516 -

```
f1_score: 0.9507 - loss: 0.1817 - val_accuracy: 0.5973 - val_f1_score:
0.5980 - val_loss: 1.2050
Epoch 14/20
12/12 _____ 8s 662ms/step - accuracy: 0.9745 -
f1_score: 0.9740 - loss: 0.1338 - val_accuracy: 0.5967 - val_f1_score:
0.5942 - val_loss: 1.4769
Epoch 15/20
12/12 _____ 8s 650ms/step - accuracy: 0.9896 -
f1_score: 0.9897 - loss: 0.0739 - val_accuracy: 0.5813 - val_f1_score:
0.5805 - val_loss: 1.5593
Epoch 16/20
12/12 _____ 7s 625ms/step - accuracy: 0.9882 -
f1_score: 0.9877 - loss: 0.0723 - val_accuracy: 0.6360 - val_f1_score:
0.6386 - val_loss: 1.2072
Epoch 17/20
12/12 _____ 7s 609ms/step - accuracy: 0.9942 -
f1_score: 0.9942 - loss: 0.0499 - val_accuracy: 0.6720 - val_f1_score:
0.6615 - val_loss: 1.1646
Epoch 18/20
12/12 _____ 7s 607ms/step - accuracy: 0.9987 -
f1_score: 0.9987 - loss: 0.0248 - val_accuracy: 0.6660 - val_f1_score:
0.6598 - val_loss: 1.2118
Epoch 19/20
12/12 _____ 7s 611ms/step - accuracy: 0.9991 -
f1_score: 0.9991 - loss: 0.0130 - val_accuracy: 0.6773 - val_f1_score:
0.6635 - val_loss: 1.2897
Epoch 20/20
12/12 _____ 7s 621ms/step - accuracy: 0.9991 -
f1_score: 0.9991 - loss: 0.0104 - val_accuracy: 0.6767 - val_f1_score:
0.6646 - val_loss: 1.2414

save_model_log('cnn_base', history_cnn_base, epochs)
plot_accuracy_f1_loss(history_cnn_base, epochs)
```





```
print(f'The best Accuracy achieved in training with cnn_base was {round(max(history_cnn_base.history['accuracy']), 4)}.')
print(f'The best F1 Score achieved in training with the cnn_base was {round(max(history_cnn_base.history['f1_score']), 4)}.')

```

The best Accuracy achieved in training with `cnn_base` was 0.9983.  
The best F1 Score achieved in training with the `cnn_base` was 0.9984.

The metrics from the `cnn_base` suggest that the training dataset has been memorized, as we observe nearly the maximum levels of accuracy and F1 score on the training set, which do not translate to the validation set, indicating its inefficiency in generalizing.

One potential cause of this overfitting could be the number of parameters, as the model contains around 22 million parameters according to the `cnn_base.summary()` function.

Considering the size of the training dataset, this may be excessive. Therefore, I will now build a model incorporating the concept of **Max Pooling**, as well as **Padding**, while also increasing the depth of the architecture.

```
num_classes = len(labels_names)

cnn_v2 = Sequential([
    layers.Input(shape=(150,150,3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),

```

```

layers.Conv2D(32, 3, padding='same', activation='relu'),
layers.MaxPooling2D(),

layers.Conv2D(64, 3, padding='same', activation='relu'),
layers.MaxPooling2D(),

layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(32, activation='relu'),
layers.Dense(num_classes)
])

cnn_v2.compile(
    optimizer='adam',
    loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
    metrics=[
        'accuracy',
        tf.keras.metrics.F1Score(average='macro', name='f1_score')
    ]
)

cnn_v2.summary()
Model: "sequential_2"

```

Layer (type) Param #	Output Shape	
conv2d_2 (Conv2D) 448	(None, 150, 150, 16)	
max_pooling2d (MaxPooling2D) 0	(None, 75, 75, 16)	
conv2d_3 (Conv2D) 4,640	(None, 75, 75, 32)	
max_pooling2d_1 (MaxPooling2D) 0	(None, 37, 37, 32)	
conv2d_4 (Conv2D) 18,496	(None, 37, 37, 64)	

0	max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 64)
0	flatten_2 (Flatten)	(None, 20736)
1,327,168	dense_6 (Dense)	(None, 64)
2,080	dense_7 (Dense)	(None, 32)
198	dense_8 (Dense)	(None, 6)
Total params: 1,353,030 (5.16 MB)		
Trainable params: 1,353,030 (5.16 MB)		
Non-trainable params: 0 (0.00 B)		

Using the `cnn_v2.summary()` function, we can clearly see the difference in the number of parameters when applying Max Pooling versus not. The model went from 22 million parameters down to just 1.3 million, while increasing its depth.

epochs = 20

```
history_cnn_v2 = cnn_v2.fit(
    train_dataset,
    validation_data=val_dataset,
    epochs=epochs
)
```

Epoch 1/20

12/12 ————— 6s 367ms/step - accuracy: 0.2053 -  
f1\_score: 0.1505 - loss: 1.8989 - val\_accuracy: 0.2813 - val\_f1\_score:  
0.2260 - val\_loss: 1.5660

Epoch 2/20

12/12 ————— 4s 341ms/step - accuracy: 0.3737 -  
f1\_score: 0.3351 - loss: 1.4589 - val\_accuracy: 0.5400 - val\_f1\_score:  
0.4904 - val\_loss: 1.2307

Epoch 3/20

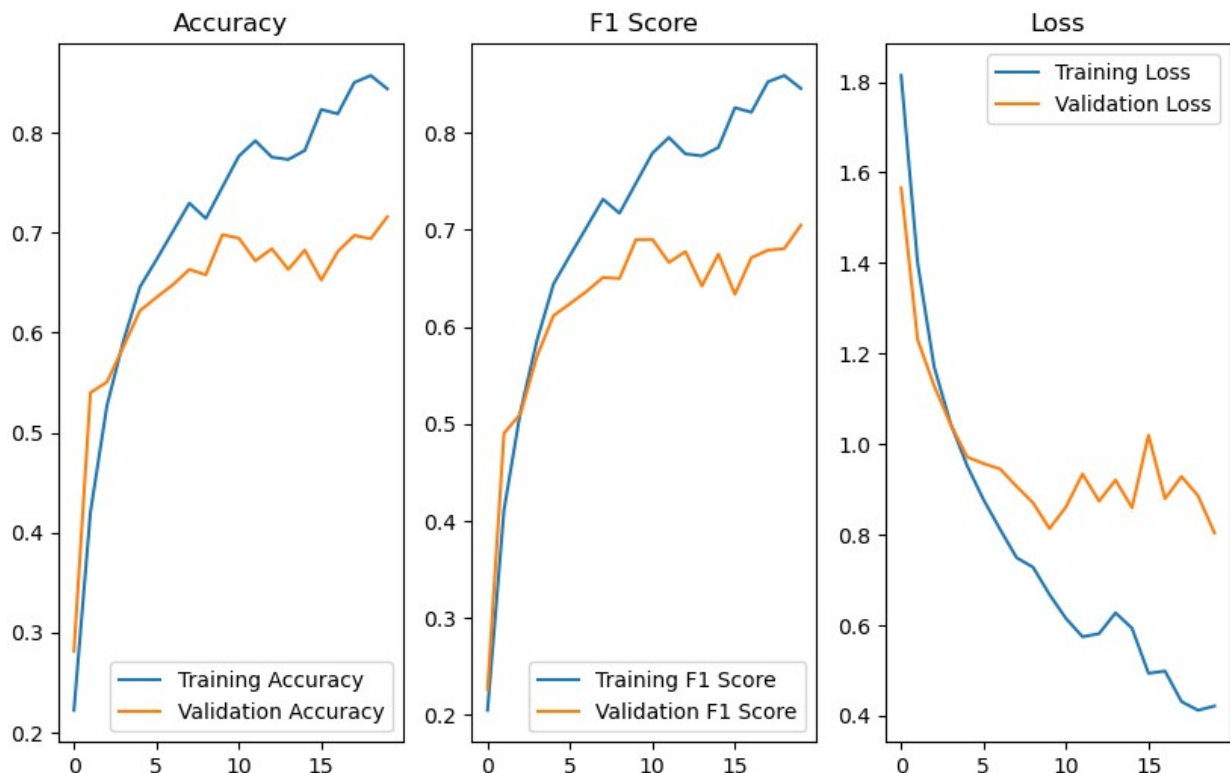
12/12 \_\_\_\_\_ 4s 321ms/step - accuracy: 0.5233 -  
f1\_score: 0.5065 - loss: 1.1891 - val\_accuracy: 0.5507 - val\_f1\_score:  
0.5104 - val\_loss: 1.1289  
Epoch 4/20  
12/12 \_\_\_\_\_ 4s 327ms/step - accuracy: 0.5816 -  
f1\_score: 0.5723 - loss: 1.0596 - val\_accuracy: 0.5860 - val\_f1\_score:  
0.5699 - val\_loss: 1.0438  
Epoch 5/20  
12/12 \_\_\_\_\_ 4s 327ms/step - accuracy: 0.6375 -  
f1\_score: 0.6365 - loss: 0.9608 - val\_accuracy: 0.6220 - val\_f1\_score:  
0.6116 - val\_loss: 0.9717  
Epoch 6/20  
12/12 \_\_\_\_\_ 4s 325ms/step - accuracy: 0.6734 -  
f1\_score: 0.6739 - loss: 0.8760 - val\_accuracy: 0.6353 - val\_f1\_score:  
0.6240 - val\_loss: 0.9570  
Epoch 7/20  
12/12 \_\_\_\_\_ 4s 320ms/step - accuracy: 0.6990 -  
f1\_score: 0.6991 - loss: 0.8222 - val\_accuracy: 0.6480 - val\_f1\_score:  
0.6366 - val\_loss: 0.9453  
Epoch 8/20  
12/12 \_\_\_\_\_ 4s 320ms/step - accuracy: 0.7320 -  
f1\_score: 0.7332 - loss: 0.7584 - val\_accuracy: 0.6633 - val\_f1\_score:  
0.6510 - val\_loss: 0.9067  
Epoch 9/20  
12/12 \_\_\_\_\_ 4s 322ms/step - accuracy: 0.7220 -  
f1\_score: 0.7238 - loss: 0.7280 - val\_accuracy: 0.6580 - val\_f1\_score:  
0.6499 - val\_loss: 0.8703  
Epoch 10/20  
12/12 \_\_\_\_\_ 4s 325ms/step - accuracy: 0.7520 -  
f1\_score: 0.7537 - loss: 0.6694 - val\_accuracy: 0.6980 - val\_f1\_score:  
0.6900 - val\_loss: 0.8133  
Epoch 11/20  
12/12 \_\_\_\_\_ 4s 323ms/step - accuracy: 0.7766 -  
f1\_score: 0.7791 - loss: 0.6238 - val\_accuracy: 0.6947 - val\_f1\_score:  
0.6902 - val\_loss: 0.8618  
Epoch 12/20  
12/12 \_\_\_\_\_ 4s 343ms/step - accuracy: 0.7989 -  
f1\_score: 0.8015 - loss: 0.5778 - val\_accuracy: 0.6720 - val\_f1\_score:  
0.6665 - val\_loss: 0.9341  
Epoch 13/20  
12/12 \_\_\_\_\_ 4s 374ms/step - accuracy: 0.7861 -  
f1\_score: 0.7881 - loss: 0.5733 - val\_accuracy: 0.6840 - val\_f1\_score:  
0.6777 - val\_loss: 0.8743  
Epoch 14/20  
12/12 \_\_\_\_\_ 4s 325ms/step - accuracy: 0.7868 -  
f1\_score: 0.7913 - loss: 0.6078 - val\_accuracy: 0.6633 - val\_f1\_score:  
0.6422 - val\_loss: 0.9206  
Epoch 15/20  
12/12 \_\_\_\_\_ 4s 328ms/step - accuracy: 0.7689 -

```

f1_score: 0.7719 - loss: 0.6085 - val_accuracy: 0.6827 - val_f1_score:
0.6751 - val_loss: 0.8598
Epoch 16/20
12/12 _____ 4s 323ms/step - accuracy: 0.8253 -
f1_score: 0.8273 - loss: 0.4933 - val_accuracy: 0.6527 - val_f1_score:
0.6337 - val_loss: 1.0198
Epoch 17/20
12/12 _____ 4s 320ms/step - accuracy: 0.8036 -
f1_score: 0.8057 - loss: 0.5301 - val_accuracy: 0.6813 - val_f1_score:
0.6714 - val_loss: 0.8795
Epoch 18/20
12/12 _____ 4s 324ms/step - accuracy: 0.8455 -
f1_score: 0.8477 - loss: 0.4422 - val_accuracy: 0.6973 - val_f1_score:
0.6791 - val_loss: 0.9285
Epoch 19/20
12/12 _____ 4s 330ms/step - accuracy: 0.8488 -
f1_score: 0.8495 - loss: 0.4262 - val_accuracy: 0.6940 - val_f1_score:
0.6808 - val_loss: 0.8863
Epoch 20/20
12/12 _____ 4s 326ms/step - accuracy: 0.8501 -
f1_score: 0.8512 - loss: 0.4138 - val_accuracy: 0.7160 - val_f1_score:
0.7050 - val_loss: 0.8040

save_model_log('cnn_v2', history_cnn_v2, epochs)
plot_accuracy_f1_loss(history_cnn_v2, epochs)

```



```
print(f'The best Accuracy achieved with cnn_v2 was  
{round(max(history_cnn_v2.history['accuracy']), 4)}'.')  
print(f'The best F1 Score achieved with the cnn_v2 was  
{round(max(history_cnn_v2.history['f1_score']), 4)}'.')
```

The best Accuracy achieved with cnn\_v2 was 0.8573.  
The best F1 Score achieved with the cnn\_v2 was 0.8592.

The divergence in metrics between the training and validation datasets suggests evidence of overfitting. However, since the validation metrics are still showing an upward trend, I will train the model for a few more epochs.

```
epochs = 10
```

```
history_cnn_v2 = cnn_v2.fit(  
    train_dataset,  
    validation_data=val_dataset,  
    epochs=epochs  
)
```

Epoch 1/10

12/12 ————— 4s 344ms/step - accuracy: 0.8334 -  
f1\_score: 0.8329 - loss: 0.4518 - val\_accuracy: 0.7027 - val\_f1\_score:  
0.6821 - val\_loss: 0.8655

Epoch 2/10

12/12 ————— 4s 339ms/step - accuracy: 0.8062 -  
f1\_score: 0.8058 - loss: 0.5319 - val\_accuracy: 0.6493 - val\_f1\_score:  
0.6315 - val\_loss: 0.9739

Epoch 3/10

12/12 ————— 4s 331ms/step - accuracy: 0.7942 -  
f1\_score: 0.7909 - loss: 0.5635 - val\_accuracy: 0.7180 - val\_f1\_score:  
0.7130 - val\_loss: 0.7780

Epoch 4/10

12/12 ————— 4s 326ms/step - accuracy: 0.8631 -  
f1\_score: 0.8644 - loss: 0.3890 - val\_accuracy: 0.7147 - val\_f1\_score:  
0.7115 - val\_loss: 0.8436

Epoch 5/10

12/12 ————— 4s 319ms/step - accuracy: 0.8735 -  
f1\_score: 0.8741 - loss: 0.3484 - val\_accuracy: 0.7153 - val\_f1\_score:  
0.7106 - val\_loss: 0.8463

Epoch 6/10

12/12 ————— 4s 325ms/step - accuracy: 0.8829 -  
f1\_score: 0.8828 - loss: 0.3229 - val\_accuracy: 0.7200 - val\_f1\_score:  
0.7114 - val\_loss: 0.8707

Epoch 7/10

12/12 ————— 4s 326ms/step - accuracy: 0.8982 -  
f1\_score: 0.8995 - loss: 0.2966 - val\_accuracy: 0.7033 - val\_f1\_score:  
0.6794 - val\_loss: 0.9020

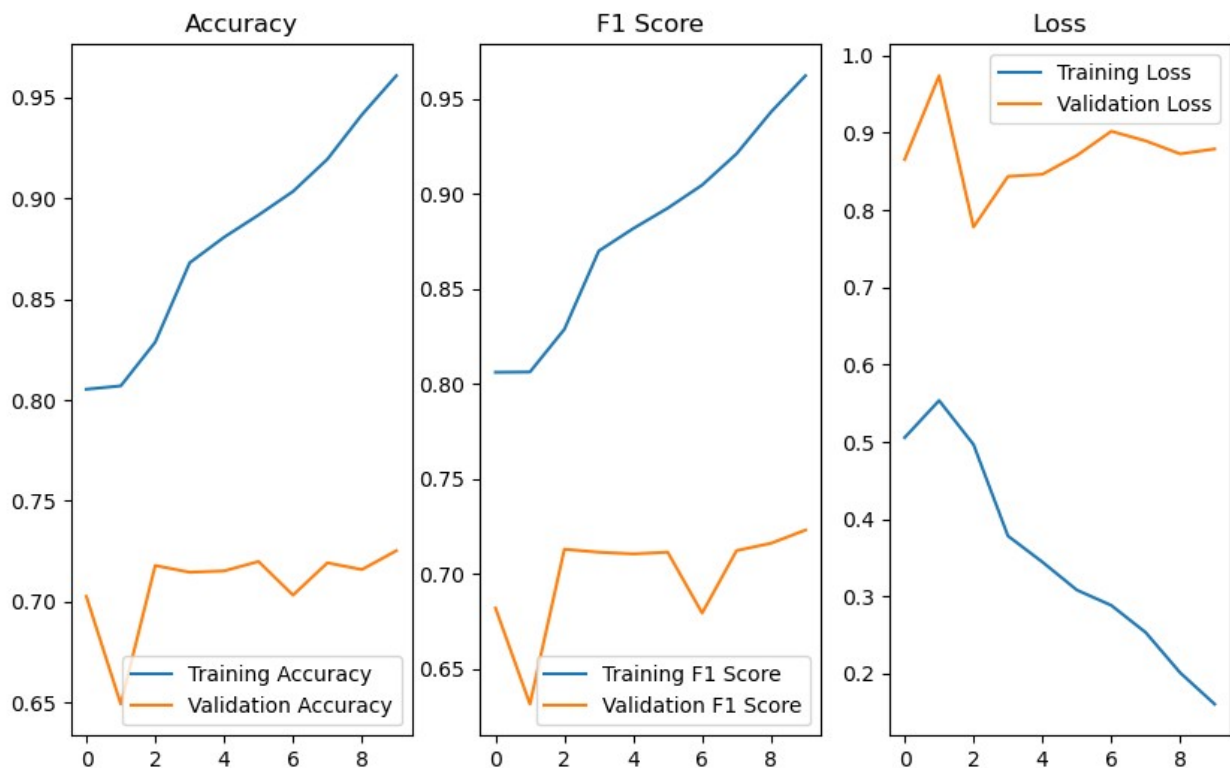
Epoch 8/10

```

12/12 ————— 4s 324ms/step - accuracy: 0.9095 -
f1_score: 0.9120 - loss: 0.2718 - val_accuracy: 0.7193 - val_f1_score:
0.7123 - val_loss: 0.8896
Epoch 9/10
12/12 ————— 4s 324ms/step - accuracy: 0.9373 -
f1_score: 0.9395 - loss: 0.2097 - val_accuracy: 0.7160 - val_f1_score:
0.7161 - val_loss: 0.8727
Epoch 10/10
12/12 ————— 4s 328ms/step - accuracy: 0.9587 -
f1_score: 0.9603 - loss: 0.1643 - val_accuracy: 0.7253 - val_f1_score:
0.7231 - val_loss: 0.8791

```

```
plot_accuracy_f1_loss(history_cnn_v2, epochs)
```



With both CNN models, `cnn_base` and `cnn_v2`, we can now observe the impact of using Max-Pooling. We transitioned from a 21M-parameter model that instantly overfitted the training data and performed poorly on the validation set, to a deeper model with only 1.5M parameters that showed slight improvements in the validation metrics — logs are available in the `log.json` file.

However, these improvements were modest (we increased accuracy from 68% to 72%), and overfitting continues to be an issue for both models.

Therefore, I will incorporate **regularization** and **dropout** techniques into the `cnn_v2` architecture to assess whether they can positively impact the model's generalization and help prevent overfitting.



```

from tensorflow.keras import regularizers

num_classes = len(labels_names)

cnn_reg = Sequential([
    layers.Input(shape=(150,150,3)),

    layers.Conv2D(16, 3, padding='same', activation='relu',
kernel_regularizer=regularizers.l2(0.01)),
    layers.MaxPooling2D(),

    layers.Conv2D(32, 3, padding='same', activation='relu',
kernel_regularizer=regularizers.l2(0.01)),
    layers.MaxPooling2D(),

    layers.Conv2D(64, 3, padding='same', activation='relu',
kernel_regularizer=regularizers.l2(0.01)),
    layers.MaxPooling2D(),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.2),
    layers.Dense(32, activation='relu'),
    layers.Dropout(0.1),
    layers.Dense(num_classes)
])

cnn_reg.compile(
    optimizer='adam',
    loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
    metrics=[
        'accuracy',
        tf.keras.metrics.F1Score(average='macro', name='f1_score')
    ]
)

```

```
cnn_reg.summary()
```

```
Model: "sequential_3"
```

Layer (type)	Output Shape
Param #	
conv2d_5 (Conv2D)	(None, 150, 150, 16)
448	
max_pooling2d_3 (MaxPooling2D)	(None, 75, 75, 16)

0				
		conv2d_6 (Conv2D)	(None, 75, 75, 32)	
4,640				
		max_pooling2d_4 (MaxPooling2D)	(None, 37, 37, 32)	
0				
		conv2d_7 (Conv2D)	(None, 37, 37, 64)	
18,496				
		max_pooling2d_5 (MaxPooling2D)	(None, 18, 18, 64)	
0				
		flatten_3 (Flatten)	(None, 20736)	
0				
		dense_9 (Dense)	(None, 64)	
1,327,168				
		dropout (Dropout)	(None, 64)	
0				
		dense_10 (Dense)	(None, 32)	
2,080				
		dropout_1 (Dropout)	(None, 32)	
0				
		dense_11 (Dense)	(None, 6)	
198				

Total params: 1,353,030 (5.16 MB)

Trainable params: 1,353,030 (5.16 MB)

Non-trainable params: 0 (0.00 B)

epochs = 25

```
history_cnn_reg = cnn_reg.fit(  
    train_dataset,  
    validation_data=val_dataset,  
    epochs=epochs  
)
```

Epoch 1/25

12/12 ————— 6s 346ms/step - accuracy: 0.2269 -  
f1\_score: 0.2037 - loss: 2.4189 - val\_accuracy: 0.4447 - val\_f1\_score:  
0.3601 - val\_loss: 1.9788

Epoch 2/25

12/12 ————— 4s 334ms/step - accuracy: 0.4117 -  
f1\_score: 0.3996 - loss: 1.9273 - val\_accuracy: 0.5553 - val\_f1\_score:  
0.4887 - val\_loss: 1.6000

Epoch 3/25

12/12 ————— 4s 325ms/step - accuracy: 0.4954 -  
f1\_score: 0.4868 - loss: 1.6749 - val\_accuracy: 0.5740 - val\_f1\_score:  
0.5128 - val\_loss: 1.4742

Epoch 4/25

12/12 ————— 4s 324ms/step - accuracy: 0.5422 -  
f1\_score: 0.5244 - loss: 1.5307 - val\_accuracy: 0.6120 - val\_f1\_score:  
0.5877 - val\_loss: 1.3583

Epoch 5/25

12/12 ————— 4s 322ms/step - accuracy: 0.5805 -  
f1\_score: 0.5704 - loss: 1.4080 - val\_accuracy: 0.6327 - val\_f1\_score:  
0.6156 - val\_loss: 1.2610

Epoch 6/25

12/12 ————— 4s 328ms/step - accuracy: 0.5972 -  
f1\_score: 0.5942 - loss: 1.3253 - val\_accuracy: 0.6247 - val\_f1\_score:  
0.5889 - val\_loss: 1.2181

Epoch 7/25

12/12 ————— 4s 323ms/step - accuracy: 0.6143 -  
f1\_score: 0.6092 - loss: 1.2802 - val\_accuracy: 0.6587 - val\_f1\_score:  
0.6403 - val\_loss: 1.1505

Epoch 8/25

12/12 ————— 4s 327ms/step - accuracy: 0.6325 -  
f1\_score: 0.6263 - loss: 1.1858 - val\_accuracy: 0.6767 - val\_f1\_score:  
0.6644 - val\_loss: 1.0898

Epoch 9/25

12/12 ————— 4s 324ms/step - accuracy: 0.6631 -  
f1\_score: 0.6636 - loss: 1.1184 - val\_accuracy: 0.6753 - val\_f1\_score:  
0.6513 - val\_loss: 1.0536

Epoch 10/25

12/12 ————— 4s 321ms/step - accuracy: 0.6726 -  
f1\_score: 0.6720 - loss: 1.1063 - val\_accuracy: 0.6867 - val\_f1\_score:  
0.6556 - val\_loss: 1.0550

Epoch 11/25

12/12 ————— 4s 322ms/step - accuracy: 0.6893 -

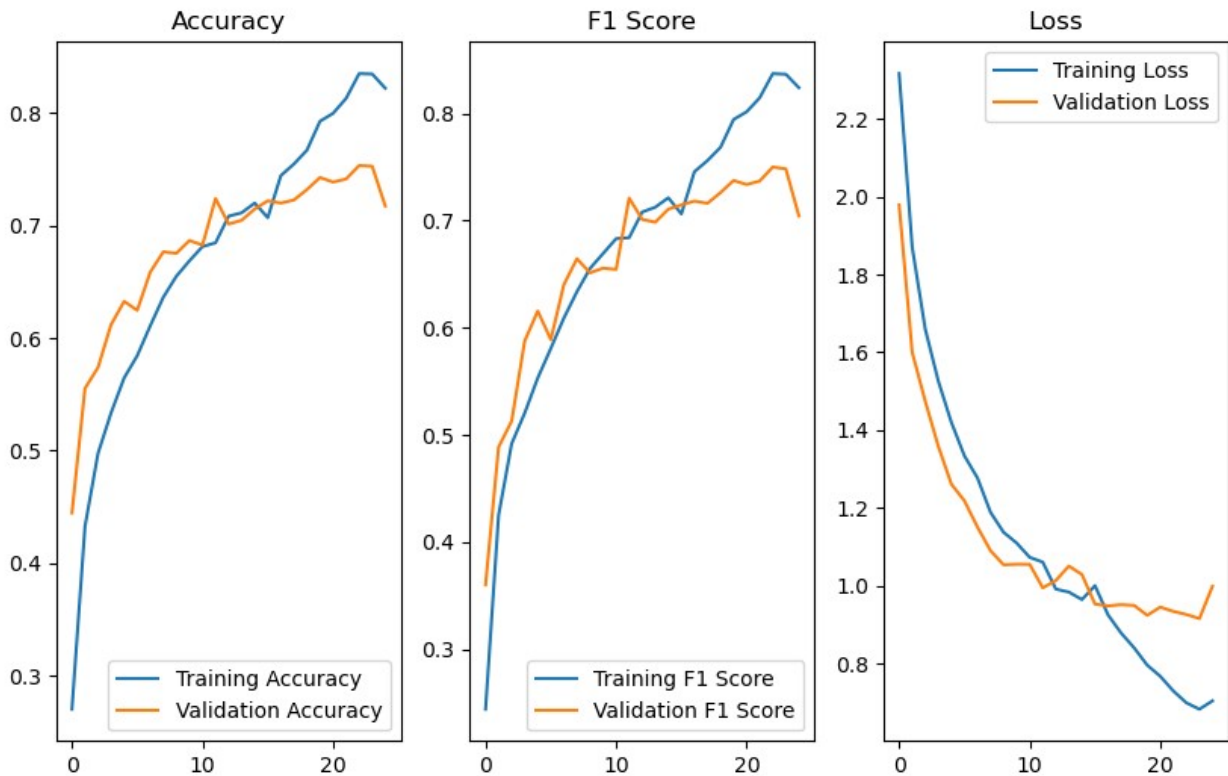
```
f1_score: 0.6890 - loss: 1.0663 - val_accuracy: 0.6827 - val_f1_score:
0.6544 - val_loss: 1.0547
Epoch 12/25
12/12 _____ 4s 329ms/step - accuracy: 0.6906 -
f1_score: 0.6847 - loss: 1.0437 - val_accuracy: 0.7240 - val_f1_score:
0.7210 - val_loss: 0.9944
Epoch 13/25
12/12 _____ 4s 327ms/step - accuracy: 0.7186 -
f1_score: 0.7181 - loss: 0.9759 - val_accuracy: 0.7013 - val_f1_score:
0.7011 - val_loss: 1.0132
Epoch 14/25
12/12 _____ 4s 327ms/step - accuracy: 0.7178 -
f1_score: 0.7198 - loss: 0.9721 - val_accuracy: 0.7047 - val_f1_score:
0.6985 - val_loss: 1.0503
Epoch 15/25
12/12 _____ 4s 325ms/step - accuracy: 0.7167 -
f1_score: 0.7171 - loss: 0.9467 - val_accuracy: 0.7147 - val_f1_score:
0.7106 - val_loss: 1.0287
Epoch 16/25
12/12 _____ 4s 326ms/step - accuracy: 0.7163 -
f1_score: 0.7175 - loss: 0.9855 - val_accuracy: 0.7220 - val_f1_score:
0.7146 - val_loss: 0.9527
Epoch 17/25
12/12 _____ 4s 325ms/step - accuracy: 0.7495 -
f1_score: 0.7511 - loss: 0.9347 - val_accuracy: 0.7200 - val_f1_score:
0.7182 - val_loss: 0.9480
Epoch 18/25
12/12 _____ 4s 327ms/step - accuracy: 0.7576 -
f1_score: 0.7597 - loss: 0.8789 - val_accuracy: 0.7227 - val_f1_score:
0.7161 - val_loss: 0.9510
Epoch 19/25
12/12 _____ 4s 327ms/step - accuracy: 0.7705 -
f1_score: 0.7727 - loss: 0.8381 - val_accuracy: 0.7320 - val_f1_score:
0.7261 - val_loss: 0.9491
Epoch 20/25
12/12 _____ 4s 331ms/step - accuracy: 0.7855 -
f1_score: 0.7868 - loss: 0.8099 - val_accuracy: 0.7427 - val_f1_score:
0.7374 - val_loss: 0.9238
Epoch 21/25
12/12 _____ 4s 331ms/step - accuracy: 0.7967 -
f1_score: 0.7989 - loss: 0.7771 - val_accuracy: 0.7387 - val_f1_score:
0.7336 - val_loss: 0.9448
Epoch 22/25
12/12 _____ 4s 331ms/step - accuracy: 0.8119 -
f1_score: 0.8145 - loss: 0.7395 - val_accuracy: 0.7413 - val_f1_score:
0.7369 - val_loss: 0.9338
Epoch 23/25
12/12 _____ 4s 329ms/step - accuracy: 0.8302 -
f1_score: 0.8330 - loss: 0.7062 - val_accuracy: 0.7533 - val_f1_score:
```

```

0.7499 - val_loss: 0.9258
Epoch 24/25
12/12 _____ 4s 328ms/step - accuracy: 0.8441 -
f1_score: 0.8467 - loss: 0.6660 - val_accuracy: 0.7527 - val_f1_score:
0.7483 - val_loss: 0.9153
Epoch 25/25
12/12 _____ 4s 326ms/step - accuracy: 0.8393 -
f1_score: 0.8421 - loss: 0.6895 - val_accuracy: 0.7173 - val_f1_score:
0.7044 - val_loss: 0.9984

plot_accuracy_f1_loss(history_cnn_reg, epochs)

```



```

epochs = 15

history_cnn_reg = cnn_reg.fit(
    train_dataset,
    validation_data=val_dataset,
    epochs=epochs
)

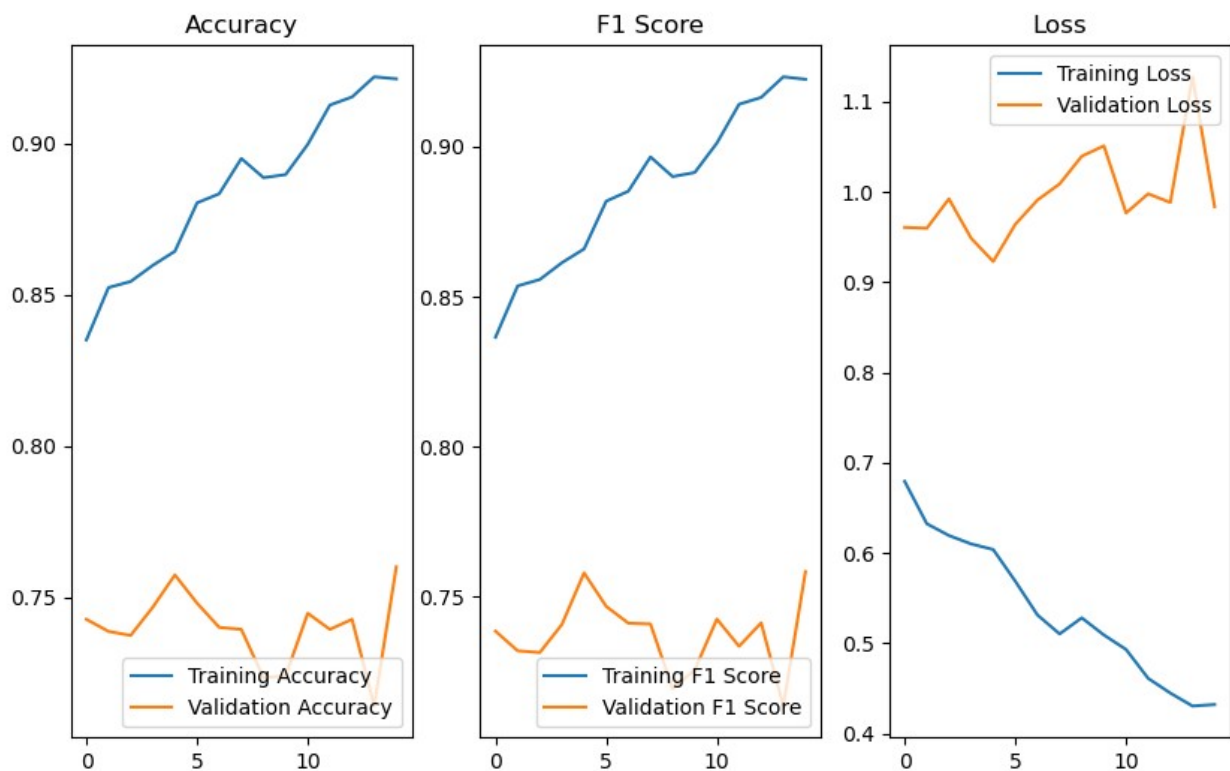
Epoch 1/15
12/12 _____ 4s 340ms/step - accuracy: 0.8287 -
f1_score: 0.8301 - loss: 0.7000 - val_accuracy: 0.7427 - val_f1_score:
0.7386 - val_loss: 0.9605
Epoch 2/15

```

12/12 \_\_\_\_\_ 4s 347ms/step - accuracy: 0.8439 -  
f1\_score: 0.8452 - loss: 0.6373 - val\_accuracy: 0.7387 - val\_f1\_score:  
0.7320 - val\_loss: 0.9595  
Epoch 3/15  
12/12 \_\_\_\_\_ 4s 329ms/step - accuracy: 0.8505 -  
f1\_score: 0.8518 - loss: 0.6297 - val\_accuracy: 0.7373 - val\_f1\_score:  
0.7315 - val\_loss: 0.9921  
Epoch 4/15  
12/12 \_\_\_\_\_ 4s 327ms/step - accuracy: 0.8579 -  
f1\_score: 0.8592 - loss: 0.6168 - val\_accuracy: 0.7467 - val\_f1\_score:  
0.7408 - val\_loss: 0.9486  
Epoch 5/15  
12/12 \_\_\_\_\_ 4s 326ms/step - accuracy: 0.8594 -  
f1\_score: 0.8613 - loss: 0.6159 - val\_accuracy: 0.7573 - val\_f1\_score:  
0.7580 - val\_loss: 0.9228  
Epoch 6/15  
12/12 \_\_\_\_\_ 4s 325ms/step - accuracy: 0.8769 -  
f1\_score: 0.8783 - loss: 0.5763 - val\_accuracy: 0.7480 - val\_f1\_score:  
0.7468 - val\_loss: 0.9640  
Epoch 7/15  
12/12 \_\_\_\_\_ 4s 372ms/step - accuracy: 0.8797 -  
f1\_score: 0.8819 - loss: 0.5361 - val\_accuracy: 0.7400 - val\_f1\_score:  
0.7412 - val\_loss: 0.9908  
Epoch 8/15  
12/12 \_\_\_\_\_ 4s 373ms/step - accuracy: 0.8905 -  
f1\_score: 0.8916 - loss: 0.5144 - val\_accuracy: 0.7393 - val\_f1\_score:  
0.7410 - val\_loss: 1.0086  
Epoch 9/15  
12/12 \_\_\_\_\_ 4s 346ms/step - accuracy: 0.8891 -  
f1\_score: 0.8904 - loss: 0.5194 - val\_accuracy: 0.7233 - val\_f1\_score:  
0.7193 - val\_loss: 1.0393  
Epoch 10/15  
12/12 \_\_\_\_\_ 4s 322ms/step - accuracy: 0.8896 -  
f1\_score: 0.8920 - loss: 0.5086 - val\_accuracy: 0.7240 - val\_f1\_score:  
0.7257 - val\_loss: 1.0508  
Epoch 11/15  
12/12 \_\_\_\_\_ 4s 326ms/step - accuracy: 0.8904 -  
f1\_score: 0.8924 - loss: 0.5063 - val\_accuracy: 0.7447 - val\_f1\_score:  
0.7426 - val\_loss: 0.9764  
Epoch 12/15  
12/12 \_\_\_\_\_ 4s 325ms/step - accuracy: 0.9130 -  
f1\_score: 0.9145 - loss: 0.4606 - val\_accuracy: 0.7393 - val\_f1\_score:  
0.7335 - val\_loss: 0.9977  
Epoch 13/15  
12/12 \_\_\_\_\_ 4s 324ms/step - accuracy: 0.9161 -  
f1\_score: 0.9167 - loss: 0.4470 - val\_accuracy: 0.7427 - val\_f1\_score:  
0.7413 - val\_loss: 0.9881  
Epoch 14/15  
12/12 \_\_\_\_\_ 4s 324ms/step - accuracy: 0.9229 -

```
f1_score: 0.9239 - loss: 0.4306 - val_accuracy: 0.7140 - val_f1_score:
0.7136 - val_loss: 1.1272
Epoch 15/15
12/12 _____ 4s 325ms/step - accuracy: 0.9152 -
f1_score: 0.9164 - loss: 0.4402 - val_accuracy: 0.7600 - val_f1_score:
0.7584 - val_loss: 0.9833

save_model_log('cnn_reg', history_cnn_reg, epochs)
plot_accuracy_f1_loss(history_cnn_reg, epochs)
```



By incorporating regularization and dropout, we observed that the model took longer to reach the point of overfitting, and the validation metrics showed a slight improvement (increasing from 72% accuracy to 76%), as seen in the log.json file.

Nevertheless, I still aim to achieve higher metrics on the validation set. Therefore, I will apply transfer learning, leveraging pre-trained models to boost accuracy and F1 score, and further improve generalization on unseen data.

## Transfer Learning

The list of the available models can be checked at the follow link:

<https://keras.io/api/applications/#available-models>.



Given the size of our dataset and the goal of iterating quickly through the training process, I will use the **MobileNet** architecture, as it is one of the lightest in terms of memory usage and number of parameters, which I believe is well-suited for our problem.

```
from tensorflow.keras.applications import MobileNet
```

```
MobileNet = MobileNet(  
    input_shape=(150,150,3),  
    include_top=False,  
    weights='imagenet',  
)
```

```
MobileNet.summary()
```

```
C:\Users\Joao\AppData\Local\Temp\ipykernel_460\4139830945.py:1:  
UserWarning: `input_shape` is undefined or non-square, or `rows` is  
not in [128, 160, 192, 224]. Weights for input shape (224, 224) will  
be loaded as the default.
```

```
MobileNet = MobileNet(  
    input_shape=(150,150,3),  
    include_top=False,  
    weights='imagenet',  
)
```

```
Model: "mobilenet_1.00_224"
```

Layer (type) Param #	Output Shape	
input_layer_4 (InputLayer) 0	(None, 150, 150, 3)	
conv1 (Conv2D) 864	(None, 75, 75, 32)	
conv1_bn (BatchNormalization) 128	(None, 75, 75, 32)	
conv1_relu (ReLU) 0	(None, 75, 75, 32)	
conv_dw_1 (DepthwiseConv2D) 288	(None, 75, 75, 32)	
conv_dw_1_bn 128	(None, 75, 75, 32)	

	(BatchNormalization)		
0	conv_dw_1_relu (ReLU)	(None, 75, 75, 32)	
2,048	conv_pw_1 (Conv2D)	(None, 75, 75, 64)	
256	conv_pw_1_bn	(None, 75, 75, 64)	
	(BatchNormalization)		
0	conv_pw_1_relu (ReLU)	(None, 75, 75, 64)	
0	conv_pad_2 (ZeroPadding2D)	(None, 76, 76, 64)	
576	conv_dw_2 (DepthwiseConv2D)	(None, 37, 37, 64)	
256	conv_dw_2_bn	(None, 37, 37, 64)	
	(BatchNormalization)		
0	conv_dw_2_relu (ReLU)	(None, 37, 37, 64)	
8,192	conv_pw_2 (Conv2D)	(None, 37, 37, 128)	
512	conv_pw_2_bn	(None, 37, 37, 128)	
	(BatchNormalization)		

0	conv_pw_2_relu (ReLU)	(None, 37, 37, 128)	
1,152	conv_dw_3 (DepthwiseConv2D)	(None, 37, 37, 128)	
512	conv_dw_3_bn (BatchNormalization)	(None, 37, 37, 128)	
0	conv_dw_3_relu (ReLU)	(None, 37, 37, 128)	
16,384	conv_pw_3 (Conv2D)	(None, 37, 37, 128)	
512	conv_pw_3_bn (BatchNormalization)	(None, 37, 37, 128)	
0	conv_pw_3_relu (ReLU)	(None, 37, 37, 128)	
0	conv_pad_4 (ZeroPadding2D)	(None, 38, 38, 128)	
1,152	conv_dw_4 (DepthwiseConv2D)	(None, 18, 18, 128)	
512	conv_dw_4_bn (BatchNormalization)	(None, 18, 18, 128)	
0	conv_dw_4_relu (ReLU)	(None, 18, 18, 128)	

conv_pw_4 (Conv2D)	(None, 18, 18, 256)	
32,768		
conv_pw_4_bn	(None, 18, 18, 256)	
1,024		
(BatchNormalization)		
conv_pw_4_relu (ReLU)	(None, 18, 18, 256)	
0		
conv_dw_5 (DepthwiseConv2D)	(None, 18, 18, 256)	
2,304		
conv_dw_5_bn	(None, 18, 18, 256)	
1,024		
(BatchNormalization)		
conv_dw_5_relu (ReLU)	(None, 18, 18, 256)	
0		
conv_pw_5 (Conv2D)	(None, 18, 18, 256)	
65,536		
conv_pw_5_bn	(None, 18, 18, 256)	
1,024		
(BatchNormalization)		
conv_pw_5_relu (ReLU)	(None, 18, 18, 256)	
0		
conv_pad_6 (ZeroPadding2D)	(None, 19, 19, 256)	
0		
conv_dw_6 (DepthwiseConv2D)	(None, 9, 9, 256)	

2,304				
		conv_dw_6_bn	(None, 9, 9, 256)	
1,024		(BatchNormalization)		
		conv_dw_6_relu (ReLU)	(None, 9, 9, 256)	
0				
		conv_pw_6 (Conv2D)	(None, 9, 9, 512)	
131,072				
		conv_pw_6_bn	(None, 9, 9, 512)	
2,048		(BatchNormalization)		
		conv_pw_6_relu (ReLU)	(None, 9, 9, 512)	
0				
		conv_dw_7 (DepthwiseConv2D)	(None, 9, 9, 512)	
4,608				
		conv_dw_7_bn	(None, 9, 9, 512)	
2,048		(BatchNormalization)		
		conv_dw_7_relu (ReLU)	(None, 9, 9, 512)	
0				
		conv_pw_7 (Conv2D)	(None, 9, 9, 512)	
262,144				
		conv_pw_7_bn	(None, 9, 9, 512)	
2,048		(BatchNormalization)		

0	conv_pw_7_relu (ReLU)	(None, 9, 9, 512)
4,608	conv_dw_8 (DepthwiseConv2D)	(None, 9, 9, 512)
2,048	conv_dw_8_bn (BatchNormalization)	(None, 9, 9, 512)
0	conv_dw_8_relu (ReLU)	(None, 9, 9, 512)
262,144	conv_pw_8 (Conv2D)	(None, 9, 9, 512)
2,048	conv_pw_8_bn (BatchNormalization)	(None, 9, 9, 512)
0	conv_pw_8_relu (ReLU)	(None, 9, 9, 512)
4,608	conv_dw_9 (DepthwiseConv2D)	(None, 9, 9, 512)
2,048	conv_dw_9_bn (BatchNormalization)	(None, 9, 9, 512)
0	conv_dw_9_relu (ReLU)	(None, 9, 9, 512)
	conv_pw_9 (Conv2D)	(None, 9, 9, 512)

262,144				
		conv_pw_9_bn	(None, 9, 9, 512)	
2,048		(BatchNormalization)		
		conv_pw_9_relu (ReLU)	(None, 9, 9, 512)	
0				
		conv_dw_10 (DepthwiseConv2D)	(None, 9, 9, 512)	
4,608				
		conv_dw_10_bn	(None, 9, 9, 512)	
2,048		(BatchNormalization)		
		conv_dw_10_relu (ReLU)	(None, 9, 9, 512)	
0				
		conv_pw_10 (Conv2D)	(None, 9, 9, 512)	
262,144				
		conv_pw_10_bn	(None, 9, 9, 512)	
2,048		(BatchNormalization)		
		conv_pw_10_relu (ReLU)	(None, 9, 9, 512)	
0				
		conv_dw_11 (DepthwiseConv2D)	(None, 9, 9, 512)	
4,608				
		conv_dw_11_bn	(None, 9, 9, 512)	
2,048		(BatchNormalization)		

0	conv_dw_11_relu (ReLU)	(None, 9, 9, 512)
262,144	conv_pw_11 (Conv2D)	(None, 9, 9, 512)
2,048	conv_pw_11_bn (BatchNormalization)	(None, 9, 9, 512)
0	conv_pw_11_relu (ReLU)	(None, 9, 9, 512)
0	conv_pad_12 (ZeroPadding2D)	(None, 10, 10, 512)
4,608	conv_dw_12 (DepthwiseConv2D)	(None, 4, 4, 512)
2,048	conv_dw_12_bn (BatchNormalization)	(None, 4, 4, 512)
0	conv_dw_12_relu (ReLU)	(None, 4, 4, 512)
524,288	conv_pw_12 (Conv2D)	(None, 4, 4, 1024)
4,096	conv_pw_12_bn (BatchNormalization)	(None, 4, 4, 1024)
	conv_pw_12_relu (ReLU)	(None, 4, 4, 1024)



0				
		conv_dw_13 (DepthwiseConv2D)	(None, 4, 4, 1024)	
9,216				
		conv_dw_13_bn	(None, 4, 4, 1024)	
4,096		(BatchNormalization)		
		conv_dw_13_relu (ReLU)	(None, 4, 4, 1024)	
0				
		conv_pw_13 (Conv2D)	(None, 4, 4, 1024)	
1,048,576				
		conv_pw_13_bn	(None, 4, 4, 1024)	
4,096		(BatchNormalization)		
		conv_pw_13_relu (ReLU)	(None, 4, 4, 1024)	
0				
Total params: 3,228,864 (12.32 MB)				
Trainable params: 3,206,976 (12.23 MB)				
Non-trainable params: 21,888 (85.50 KB)				

As I intend to use the pre-trained and already optimized weights, I will freeze these parameters.

```
for layer in MobileNet.layers:
    layer.trainable = False

# To check that there are no Trainable params now:
MobileNet.summary()
Model: "mobilenet_1.00_224"
```

Layer (type) Param #	Output Shape	
input_layer_4 (InputLayer) 0	(None, 150, 150, 3)	
conv1 (Conv2D) 864	(None, 75, 75, 32)	
conv1_bn (BatchNormalization) 128	(None, 75, 75, 32)	
conv1_relu (ReLU) 0	(None, 75, 75, 32)	
conv_dw_1 (DepthwiseConv2D) 288	(None, 75, 75, 32)	
conv_dw_1_bn (BatchNormalization) 128	(None, 75, 75, 32)	
conv_dw_1_relu (ReLU) 0	(None, 75, 75, 32)	
conv_pw_1 (Conv2D) 2,048	(None, 75, 75, 64)	
conv_pw_1_bn (BatchNormalization) 256	(None, 75, 75, 64)	
conv_pw_1_relu (ReLU) 0	(None, 75, 75, 64)	

0	conv_pad_2 (ZeroPadding2D)	(None, 76, 76, 64)	
576	conv_dw_2 (DepthwiseConv2D)	(None, 37, 37, 64)	
256	conv_dw_2_bn (BatchNormalization)	(None, 37, 37, 64)	
0	conv_dw_2_relu (ReLU)	(None, 37, 37, 64)	
8,192	conv_pw_2 (Conv2D)	(None, 37, 37, 128)	
512	conv_pw_2_bn (BatchNormalization)	(None, 37, 37, 128)	
0	conv_pw_2_relu (ReLU)	(None, 37, 37, 128)	
1,152	conv_dw_3 (DepthwiseConv2D)	(None, 37, 37, 128)	
512	conv_dw_3_bn (BatchNormalization)	(None, 37, 37, 128)	
0	conv_dw_3_relu (ReLU)	(None, 37, 37, 128)	
16,384	conv_pw_3 (Conv2D)	(None, 37, 37, 128)	

512	conv_pw_3_bn	(None, 37, 37, 128)	
	(BatchNormalization)		
0	conv_pw_3_relu (ReLU)	(None, 37, 37, 128)	
0	conv_pad_4 (ZeroPadding2D)	(None, 38, 38, 128)	
1,152	conv_dw_4 (DepthwiseConv2D)	(None, 18, 18, 128)	
512	conv_dw_4_bn	(None, 18, 18, 128)	
	(BatchNormalization)		
0	conv_dw_4_relu (ReLU)	(None, 18, 18, 128)	
32,768	conv_pw_4 (Conv2D)	(None, 18, 18, 256)	
1,024	conv_pw_4_bn	(None, 18, 18, 256)	
	(BatchNormalization)		
0	conv_pw_4_relu (ReLU)	(None, 18, 18, 256)	
2,304	conv_dw_5 (DepthwiseConv2D)	(None, 18, 18, 256)	
1,024	conv_dw_5_bn	(None, 18, 18, 256)	
	(BatchNormalization)		

0	conv_dw_5_relu (ReLU)	(None, 18, 18, 256)
65,536	conv_pw_5 (Conv2D)	(None, 18, 18, 256)
1,024	conv_pw_5_bn (BatchNormalization)	(None, 18, 18, 256)
0	conv_pw_5_relu (ReLU)	(None, 18, 18, 256)
0	conv_pad_6 (ZeroPadding2D)	(None, 19, 19, 256)
2,304	conv_dw_6 (DepthwiseConv2D)	(None, 9, 9, 256)
1,024	conv_dw_6_bn (BatchNormalization)	(None, 9, 9, 256)
0	conv_dw_6_relu (ReLU)	(None, 9, 9, 256)
131,072	conv_pw_6 (Conv2D)	(None, 9, 9, 512)
2,048	conv_pw_6_bn (BatchNormalization)	(None, 9, 9, 512)
0	conv_pw_6_relu (ReLU)	(None, 9, 9, 512)

4,608	conv_dw_7 (DepthwiseConv2D)	(None, 9, 9, 512)
2,048	conv_dw_7_bn	(None, 9, 9, 512)
	(BatchNormalization)	
0	conv_dw_7_relu (ReLU)	(None, 9, 9, 512)
262,144	conv_pw_7 (Conv2D)	(None, 9, 9, 512)
2,048	conv_pw_7_bn	(None, 9, 9, 512)
	(BatchNormalization)	
0	conv_pw_7_relu (ReLU)	(None, 9, 9, 512)
4,608	conv_dw_8 (DepthwiseConv2D)	(None, 9, 9, 512)
2,048	conv_dw_8_bn	(None, 9, 9, 512)
	(BatchNormalization)	
0	conv_dw_8_relu (ReLU)	(None, 9, 9, 512)
262,144	conv_pw_8 (Conv2D)	(None, 9, 9, 512)
2,048	conv_pw_8_bn	(None, 9, 9, 512)

	(BatchNormalization)		
0	conv_pw_8_relu (ReLU)	(None, 9, 9, 512)	
4,608	conv_dw_9 (DepthwiseConv2D)	(None, 9, 9, 512)	
2,048	conv_dw_9_bn	(None, 9, 9, 512)	
	(BatchNormalization)		
0	conv_dw_9_relu (ReLU)	(None, 9, 9, 512)	
262,144	conv_pw_9 (Conv2D)	(None, 9, 9, 512)	
2,048	conv_pw_9_bn	(None, 9, 9, 512)	
	(BatchNormalization)		
0	conv_pw_9_relu (ReLU)	(None, 9, 9, 512)	
4,608	conv_dw_10 (DepthwiseConv2D)	(None, 9, 9, 512)	
2,048	conv_dw_10_bn	(None, 9, 9, 512)	
	(BatchNormalization)		
0	conv_dw_10_relu (ReLU)	(None, 9, 9, 512)	

conv_pw_10 (Conv2D)	(None, 9, 9, 512)	
262,144		
conv_pw_10_bn	(None, 9, 9, 512)	
2,048		
(BatchNormalization)		
conv_pw_10_relu (ReLU)	(None, 9, 9, 512)	
0		
conv_dw_11 (DepthwiseConv2D)	(None, 9, 9, 512)	
4,608		
conv_dw_11_bn	(None, 9, 9, 512)	
2,048		
(BatchNormalization)		
conv_dw_11_relu (ReLU)	(None, 9, 9, 512)	
0		
conv_pw_11 (Conv2D)	(None, 9, 9, 512)	
262,144		
conv_pw_11_bn	(None, 9, 9, 512)	
2,048		
(BatchNormalization)		
conv_pw_11_relu (ReLU)	(None, 9, 9, 512)	
0		
conv_pad_12 (ZeroPadding2D)	(None, 10, 10, 512)	
0		
conv_dw_12 (DepthwiseConv2D)	(None, 4, 4, 512)	
4,608		



2,048	conv_dw_12_bn (BatchNormalization)	(None, 4, 4, 512)	
0	conv_dw_12_relu (ReLU)	(None, 4, 4, 512)	
524,288	conv_pw_12 (Conv2D)	(None, 4, 4, 1024)	
4,096	conv_pw_12_bn (BatchNormalization)	(None, 4, 4, 1024)	
0	conv_pw_12_relu (ReLU)	(None, 4, 4, 1024)	
9,216	conv_dw_13 (DepthwiseConv2D)	(None, 4, 4, 1024)	
4,096	conv_dw_13_bn (BatchNormalization)	(None, 4, 4, 1024)	
0	conv_dw_13_relu (ReLU)	(None, 4, 4, 1024)	
1,048,576	conv_pw_13 (Conv2D)	(None, 4, 4, 1024)	
4,096	conv_pw_13_bn (BatchNormalization)	(None, 4, 4, 1024)	
0	conv_pw_13_relu (ReLU)	(None, 4, 4, 1024)	

```
Total params: 3,228,864 (12.32 MB)
Trainable params: 0 (0.00 B)
Non-trainable params: 3,228,864 (12.32 MB)
```

```
num_classes = len(labels_names)

cnn_mobilenet = Sequential([
    MobileNet,
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(num_classes)
])

cnn_mobilenet.summary()
Model: "sequential_4"
```

Layer (type)	Output Shape
Param #	
mobilenet_1.00_224 (Functional)	(None, 4, 4, 1024)
3,228,864	
flatten_4 (Flatten)	(None, 16384)
0	
dense_12 (Dense)	(None, 64)
1,048,640	
dropout_2 (Dropout)	(None, 64)
0	
dense_13 (Dense)	(None, 6)
390	

```
Total params: 4,277,894 (16.32 MB)
```

Trainable params: 1,049,030 (4.00 MB)

Non-trainable params: 3,228,864 (12.32 MB)

```
cnn_mobilenet.compile(  
    optimizer='adam',  
    loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),  
    metrics=[  
        'accuracy',  
        tf.keras.metrics.F1Score(average='macro', name='f1_score')  
    ]  
)
```

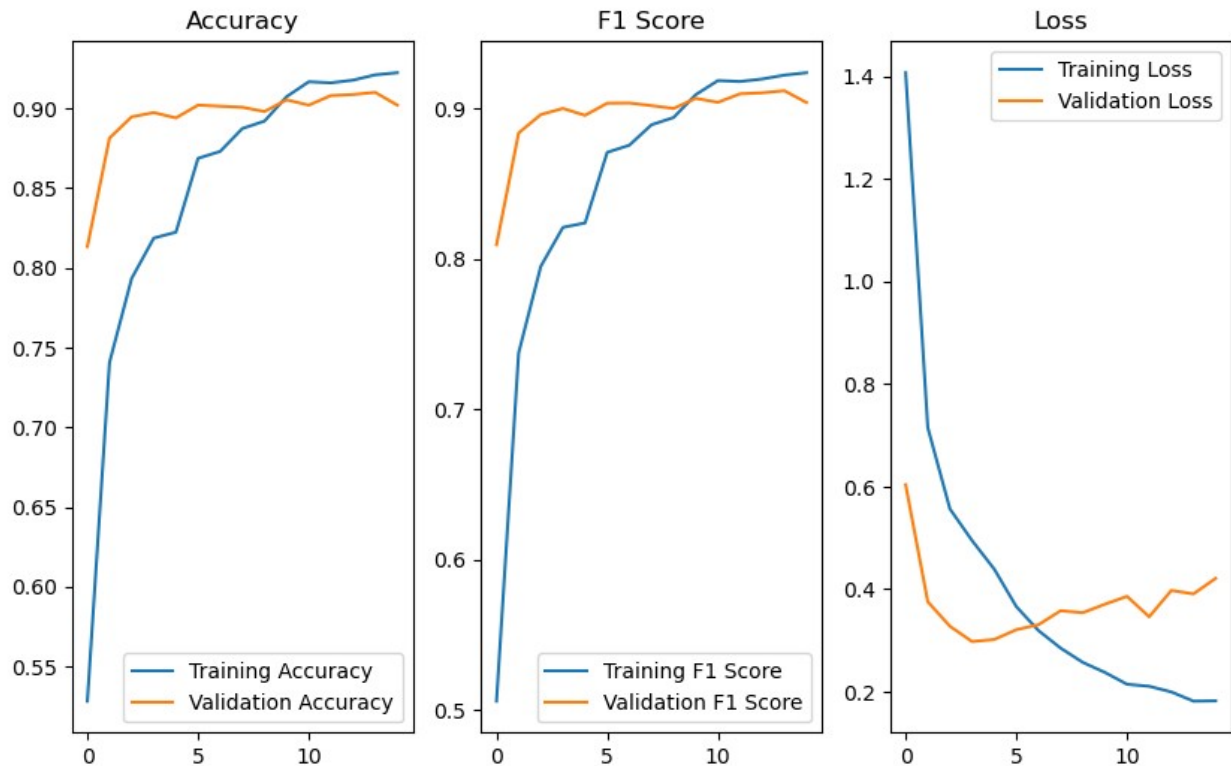
Since I expect this to be the best-performing model among all those created so far, I will save its structure (both architecture and weights) from the best-performing epoch — specifically, the one that achieved the highest accuracy on the validation set.

To do this, I will use the `ModelCheckpoint()` class, as described at the following link: [https://keras.io/api/callbacks/model\\_checkpoint/](https://keras.io/api/callbacks/model_checkpoint/), and incorporate it into the `.fit()` method.

```
from tensorflow.keras.callbacks import ModelCheckpoint  
  
checkpoint = ModelCheckpoint(  
    filepath='cnn_mobilenet.keras',  
    monitor='val_accuracy',  
    mode='max',  
    save_best_only=True,  
    save_weights_only=False  
)  
  
epochs = 15  
  
history_mobilenet = cnn_mobilenet.fit(  
    train_dataset,  
    validation_data=val_dataset,  
    epochs=epochs,  
    callbacks=[checkpoint]  
)  
  
Epoch 1/15  
12/12 _____ 12s 870ms/step - accuracy: 0.4191 -  
f1_score: 0.3931 - loss: 2.0674 - val_accuracy: 0.8133 - val_f1_score:  
0.8092 - val_loss: 0.6037  
Epoch 2/15  
12/12 _____ 11s 900ms/step - accuracy: 0.7452 -  
f1_score: 0.7369 - loss: 0.7112 - val_accuracy: 0.8813 - val_f1_score:  
0.8836 - val_loss: 0.3756  
Epoch 3/15  
12/12 _____ 11s 949ms/step - accuracy: 0.7896 -  
f1_score: 0.7909 - loss: 0.5599 - val_accuracy: 0.8947 - val_f1_score:
```

```
0.8958 - val_loss: 0.3277
Epoch 4/15
12/12 _____ 9s 790ms/step - accuracy: 0.8151 -
f1_score: 0.8171 - loss: 0.4955 - val_accuracy: 0.8973 - val_f1_score:
0.8999 - val_loss: 0.2980
Epoch 5/15
12/12 _____ 10s 882ms/step - accuracy: 0.8288 -
f1_score: 0.8294 - loss: 0.4275 - val_accuracy: 0.8940 - val_f1_score:
0.8953 - val_loss: 0.3019
Epoch 6/15
12/12 _____ 11s 939ms/step - accuracy: 0.8749 -
f1_score: 0.8769 - loss: 0.3563 - val_accuracy: 0.9020 - val_f1_score:
0.9033 - val_loss: 0.3211
Epoch 7/15
12/12 _____ 9s 765ms/step - accuracy: 0.8734 -
f1_score: 0.8757 - loss: 0.3212 - val_accuracy: 0.9013 - val_f1_score:
0.9035 - val_loss: 0.3307
Epoch 8/15
12/12 _____ 10s 858ms/step - accuracy: 0.8864 -
f1_score: 0.8874 - loss: 0.2932 - val_accuracy: 0.9007 - val_f1_score:
0.9017 - val_loss: 0.3581
Epoch 9/15
12/12 _____ 10s 869ms/step - accuracy: 0.8842 -
f1_score: 0.8843 - loss: 0.2602 - val_accuracy: 0.8980 - val_f1_score:
0.8999 - val_loss: 0.3542
Epoch 10/15
12/12 _____ 10s 849ms/step - accuracy: 0.9030 -
f1_score: 0.9044 - loss: 0.2472 - val_accuracy: 0.9053 - val_f1_score:
0.9065 - val_loss: 0.3708
Epoch 11/15
12/12 _____ 10s 883ms/step - accuracy: 0.9200 -
f1_score: 0.9212 - loss: 0.2001 - val_accuracy: 0.9020 - val_f1_score:
0.9039 - val_loss: 0.3861
Epoch 12/15
12/12 _____ 11s 906ms/step - accuracy: 0.9152 -
f1_score: 0.9155 - loss: 0.2156 - val_accuracy: 0.9080 - val_f1_score:
0.9096 - val_loss: 0.3462
Epoch 13/15
12/12 _____ 10s 890ms/step - accuracy: 0.9163 -
f1_score: 0.9172 - loss: 0.2043 - val_accuracy: 0.9087 - val_f1_score:
0.9102 - val_loss: 0.3976
Epoch 14/15
12/12 _____ 11s 940ms/step - accuracy: 0.9242 -
f1_score: 0.9242 - loss: 0.1799 - val_accuracy: 0.9100 - val_f1_score:
0.9116 - val_loss: 0.3909
Epoch 15/15
12/12 _____ 11s 947ms/step - accuracy: 0.9264 -
f1_score: 0.9272 - loss: 0.1677 - val_accuracy: 0.9020 - val_f1_score:
0.9038 - val_loss: 0.4212
```

```
save_model_log('cnn_mobilenet', history_mobilenet, epochs)
plot_accuracy_f1_loss(history_mobilenet, epochs)
```



Although the training metrics are still on an upwards trend, in the last few epochs the training and validation metrics have started to diverge, so I've decided not to continue training this architecture.

Even so, this clearly demonstrates the strength of using transfer learning instead of building a neural network from scratch, as the validation metrics were already better, even before any training, than those of all the previously tested models.

That said, since this is our best-performing model so far, I will use the saved iteration to evaluate its performance on the test dataset.

## Assessing the best model on the test dataset

Now that the best-performing iteration based on validation accuracy has been saved to the filepath **cnn\_mobilenet.keras**, I will load this model using the `load_model()` class and evaluate its performance on the test dataset.

```
from tensorflow.keras.models import load_model
best_cnn = load_model('cnn_mobilenet.keras')
test_loss, test_accuracy, test_f1_score =
```

```
best_cnn.evaluate(test_dataset)

print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}%")
print(f"Test F1 Score: {test_f1_score:.4f}")

3/3 _____ 2s 370ms/step - accuracy: 0.9183 - f1_score:
0.9239 - loss: 0.4003
Test Loss: 0.4086
Test Accuracy: 0.9167%
Test F1 Score: 0.9214
```

Since the training methodology has been validated and the model has achieved strong performance across the training, validation, and test datasets, I will now proceed to retrain this architecture using the full dataset in order to maximize the data available for final model training.

## Final Training

We still have the entire dataset stored in the **dataset** variable, which has already been shuffled. However, I will need to normalize the data, split it into batches, and apply `prefetch(tf.data.AUTOTUNE)`: This is a best practice in data pipelines, especially when training with large datasets, as it helps optimize performance by overlapping data loading and model training.

```
len(dataset)

14034

dataset = dataset.map(normalize_image)
dataset = dataset.batch(32)
dataset = dataset.prefetch(tf.data.AUTOTUNE)

save_best = ModelCheckpoint(
    filepath='final_cnn.keras',
    monitor='accuracy',
    mode='max',
    save_best_only=True,
    save_weights_only=False
)

best_cnn.fit(
    dataset,
    epochs=10,
    callbacks=[save_best]
)
```

```
Epoch 1/10
439/439 _____ 40s 90ms/step - accuracy: 0.9374 -
f1_score: 0.9394 - loss: 0.1650
Epoch 2/10
439/439 _____ 44s 100ms/step - accuracy: 0.9454 -
f1_score: 0.9472 - loss: 0.1332
Epoch 3/10
439/439 _____ 44s 100ms/step - accuracy: 0.9414 -
f1_score: 0.9433 - loss: 0.1450
Epoch 4/10
439/439 _____ 44s 99ms/step - accuracy: 0.9443 -
f1_score: 0.9459 - loss: 0.1397
Epoch 5/10
439/439 _____ 44s 100ms/step - accuracy: 0.9451 -
f1_score: 0.9468 - loss: 0.1431
Epoch 6/10
439/439 _____ 44s 99ms/step - accuracy: 0.9427 -
f1_score: 0.9446 - loss: 0.1393
Epoch 7/10
439/439 _____ 43s 98ms/step - accuracy: 0.9447 -
f1_score: 0.9462 - loss: 0.1454
Epoch 8/10
439/439 _____ 43s 98ms/step - accuracy: 0.9490 -
f1_score: 0.9505 - loss: 0.1279
Epoch 9/10
439/439 _____ 38s 87ms/step - accuracy: 0.9466 -
f1_score: 0.9482 - loss: 0.1351
Epoch 10/10
439/439 _____ 35s 79ms/step - accuracy: 0.9464 -
f1_score: 0.9481 - loss: 0.1494

<keras.src.callbacks.history.History at 0x239cf0ebd70>
```

The model was trained above for a total of 30 epochs (10+10+10).

During the last 10 epochs, all metrics began to exhibit a bit of a ping-pong effect, indicating that the model was struggling to improve consistently. As a result, I decided to stop the training.

The best iteration in terms of accuracy was saved in the file `final_cnn.keras`, achieving **94% accuracy**. This represents the final performance for the problem, and the model would now be ready to be evaluated on new datasets and further improved.